

## 02 Optimisation

### Optimisers

#### modern optimisers

##### from SGD to Adam

###### SGD (1951)

$$\theta = \theta - \alpha \cdot \nabla L$$

Simple but slow, sensitive to learning rate

###### Momentum (1964)

$$v = \beta \cdot v + \nabla L$$
$$\theta = \theta - \eta \cdot v$$

- computes gradients at current point
- momentum is a coarse prediction of where SGD will move

###### Nesterov accelerated momentum

$$m_{t+1} = \beta \cdot m_t + (1 - \beta) \sum_i \in \mathcal{B}_t \frac{\partial \ell_i [w_t - \eta \cdot m_t]}{\partial w}$$
$$w_{t+1} = w_t - \eta \cdot m_{t+1}$$

- Nesterov accelerated momentum computes the gradient at a predicted point of where the gradient **will** move after the modification
- frequently, though not always, better than standard momentum method

###### RMSprop (2012)

$$s = \beta \cdot s + (1 - \beta) \cdot \nabla L^2$$
$$\theta = \theta - \alpha \cdot \nabla L / \sqrt{s}$$

- Adapts learning rate per parameter

###### Adam (2014)

$$\begin{array}{ll} m = \beta_1 \cdot m + (1 - \beta_1) \cdot \nabla L & \text{Momentum} \\ v = \beta_2 \cdot v + (1 - \beta_2) \cdot \nabla L^2 & \text{RMSprop} \\ \theta = \theta - \eta \cdot \hat{m} / \sqrt{\hat{v}} & \text{Bias-corrected} \end{array}$$

- Best of both worlds!
  - takes into account the variance
- 

## Why Adam Works

### Per-Parameter Learning Rates

**Problem:** Different parameters need different learning rates

- Embeddings: Large gradients, need small LR
- Output layer: Small gradients, need large LR
- vanishing gradient problem too

**Adam's Solution:**

$$\text{Effective LR} = \alpha / \sqrt{(\text{variance} + \epsilon)}$$

- High variance  $\rightarrow$  Lower LR (stable)
  - Low variance  $\rightarrow$  Higher LR (faster)
- 

## AdamW - the important fix

### problem with Adam + L2

**What we want:**

$$\theta = \theta - \alpha(\nabla L + \lambda\theta) \quad \text{Decoupled weight decay}$$

**What Adam + L2 does:**

$$m = \beta_1 m + (1 - \beta_1)(\nabla L + \lambda\theta)$$

wrong couples decay with adaptive LR

### AdamW fix

Apply weight decay AFTER Adam update

$$\theta = \theta - \alpha \cdot \hat{m} / \sqrt{\hat{v}} - \lambda \theta$$

- Transformers always use AdamW
  - Better generalisation
  - Proper regularisation
- 

## Learning Rate Schedules

### warmup (Critical for Transformers)

```
if step < warmup_steps:  
    lr = base_lr * step / warmup_steps
```

Why: Large initial gradients can destabilise

### Cosine Annealing

```
lr = min_lr + (max_lr - min_lr) * 0.5 * (1 + cos(pi * step / total_steps))
```

Why: Smooth decay, allows "restarts"

### Inverse Square Root (GPT-style)

```
lr = base_lr * min(step^(-0.5), step * warmup_steps^(-1.5))
```

Why: Fast initial progress, slow refinement

---

---

## generalisation

### the generalisation puzzle

### the mystery

#### What we observe:

- Models with 100B+ parameters training on "only" millions of examples
- Perfect memorisation of training data
- Yet they generalise!

**Classical theory says:** This shouldn't work!

$$\epsilon^2 < \frac{\log |H_\epsilon| + \log 1/\delta}{2m}$$

- where
  - $\epsilon$  is the difference between training and generalisation error
  - $|H_\epsilon|$  is the  $\epsilon$  cover of the hypothesis class, usually  $|H_\epsilon| \sim (1/\epsilon)^d$
  - $d$  is the VC dimension
  - $\delta$  is confidence
  - $m$  is the number of samples
- VC dimension  $\rightarrow \infty$  for deep networks
- Rademacher complexity  $\rightarrow \infty$
- PAC bounds  $\rightarrow$  vacuous

## the resolution:

- SGD has implicit bias
- Networks prefer simple functions
- Real data has structure

---

## implicit regularisation

### SGD finds special solutions

All these solutions have zero training loss:

```
θ₁ = [1000, -999, 0.1, ...] # Complex
θ₂ = [0.1, 0.1, 0.1, ...]   # Simple
```

SGD prefers  $\theta_2$ !

### mechanisms:

1. **Noise**: SGD noise prevents overfitting
2. **Early stopping**: Implicit regularization
3. **Architecture**: CNNs bias toward local features
4. **Initialization**: Start near simple functions

**Net effect:** Loss = Training\_Loss + Implicit\_Complexity\_Penalty

---

# grokking - delayed generalization

## the phenomenon

Epoch 1–1000: Train = 100%, Test = 50% (memorization)

Epoch 1001–5000: Train = 100%, Test = 50% (still memorizing)

Epoch 5001: Train = 100%, Test = 99% (sudden generalization!)

## why grokking happens:

1. **Phase 1:** Memorize with complex pattern
2. **Phase 2:** Simplify internal representation
3. **Phase 3:** Simple pattern generalizes

**Example:** Modular arithmetic

- Memorize: Store all input-output pairs
- Generalize: Learn the algorithm

**Lesson:** Generalization can emerge long after fitting!

---

## what really drives generalization?

### the modern view

**Not (just) explicit regularization:**

- L2, dropout help but aren't necessary
- Unregularized networks still generalize

**But rather:**

1. **Inductive bias** from architecture
2. **Implicit bias** from optimization
3. **Data structure** (manifold hypothesis)
4. **Overparameterization** (lottery tickets)

## the simplicity bias

Networks prefer simple functions that:

- Vary smoothly
- Have low frequency
- Decompose hierarchically

- Match data symmetries
- networks are lazy and want to compute as little as possible

 **Key:** Deep learning works because real-world data is simple in the right coordinate system!

---




---

## double descent curve

### classical vs modern wisdom

#### classical U-curve (a bias-variance tradeoff)



**Belief:** Optimal complexity balances bias and variance

### the bias-variance tradeoff

$$\begin{aligned} Err(\hat{f}) &= E[(y - \hat{f})^2] \\ &= E[y^2] - 2E[y\hat{f}] + E[\hat{f}^2] \end{aligned}$$

where

$$\begin{aligned} E[y^2] &= E[(f + \epsilon)^2] = E[f^2] + 2E[f\epsilon] + \epsilon^2 \\ &= E[f^2] + 2E[f] \cdot 0 + \epsilon^2 \\ &= f^2 + \sigma^2 \end{aligned}$$

where

$$Var(X) = E[(X - E[X])^2] = E[X^2] - E[X]^2$$


thus

$$\begin{aligned} E[\hat{f}^2] &= Var(\hat{f}) + E[\hat{f}]^2 \\ E[y\hat{f}] &= E[(f + \epsilon)\hat{f}] \\ &= E[f\hat{f}] + E[\epsilon\hat{f}] \\ &= fE[\hat{f}] \end{aligned}$$

thus

$$\begin{aligned} Err[\hat{f}] &= f^2 + \sigma^2 - 2fE[\hat{f}] + Var(\hat{f}) + E[\hat{f}]^2 \\ &= f^2 - 2fE[\hat{f}] + E[\hat{f}]^2 + \sigma^2 + Var(\hat{f}) \\ &= (f - E[\hat{f}])^2 + \sigma^2 + Var(\hat{f}) \\ &= Bias[\hat{f}]^2 + \sigma^2 + Var(\hat{f}) \end{aligned}$$


- **bias** is the average error between average prediction and the true values
- **variance** (a system error) is the residual noisy error - the differences between average predictions
- minimising one may lead to the growth of the other (see drawing below)
  - multiple expert model may be a solution
- the system noise can be further broken down to stable pattern noise, occasion noise, and level noise






[MLU-EXPLAIN: the bias-variance tradeoff](#)

---

## modern double descent



in reality we can keep going past interpolation!



(after [Reconciling modern machine learning practice and bias-variance trade-off](#), Belkin et al. 2018) Actually not very easy to copy in code

[MLU-EXPLAIN: double descent](#)

[MLU-EXPLAIN: double descent simple math](#)

## the interpolation threshold

### three regimes

#### 1. under-parameterised ( $n_{\text{params}} < n_{\text{data}}$ )

- can't fit all training data
- classical bias-variance tradeoff
- test error follows U-curve

#### 2. critical ( $n_{\text{params}} \approx n_{\text{data}}$ )

- just enough to memorize
- worst generalisation!

- peak test error

### 3. over-parameterised ( $n_{\text{params}} \gg n_{\text{data}}$ )

- many ways to fit data
- SGD chooses simple solution
- test error decreases again!

**at interpolation threshold, only ONE solution fits → It's complex!**

---

## why double descent happens

### mathematical intuition

**underparameterized:** Unique solution via least squares

$$\theta^* = (X^T X)^{(-1)} X^T y$$

**critical regime:** matrix barely invertible

$$\theta^* = (X^T X + \varepsilon I)^{(-1)} X^T y \quad \# \varepsilon \rightarrow 0, \|\theta^*\| \rightarrow \infty$$

**Overparameterized:** Minimum norm solution

$$\theta^* = X^+ y = X^T (X X^T)^{(-1)} y \quad \# \text{Well-conditioned!}$$

**key line: more parameters → more solutions → SGD picks simpler one**

---

## practical implications

### do's:

- use large models when possible
- train past interpolation
- do not stop at 100% training accuracy
- use early stopping on validation loss, **not train**

### don'ts:

- don't choose model size to match data size
- don't fear overparameterization
- don't trust classical statistical guidance

## where double descent appears:

- Model size (width, depth)
- Training time (epoch-wise)
- Data size (sample-wise)
- Even label noise!

**modern practice: go big and regularise implicitly via SGD**

---

---

## loss function visualisation

### the geometry of optimisation

### high-dimensional reality

**the curse of visualisation:**

- networks have millions of dimensions
- we can only see 2D/3D projections
- most intuition from low-dimensions is WRONG

**what is different in high-dimensionality**

- no local minima (mostly saddles)
- volume at surface (not center)
- all points equidistant
- connected solution manifolds

---

## visualisation techniques

### 1. random projections

```
θ_proj = θ* + α·δ₁ + β·δ₂ # δᵢ ~ N(0,1)
```

Shows local geometry around minimum

### 2. PCA projections

```
θ_proj = θ* + α·PC₁ + β·PC₂
```

Shows dominant variation directions

### 3. linear interpolation

$$\theta_{\text{interp}} = (1-\alpha) \cdot \theta_1 + \alpha \cdot \theta_2$$

Tests mode connectivity

### 4. filter normalisation (Li et al. 2018)


$$\delta_{\text{normalized}} = \delta * ||\theta|| / ||\delta|| \quad \# \text{ Scale-invariant}$$

Handles different parameter scales

---


## sharp vs flat minima

**sharp minimum:**



poor generalisation

**flat minimum:**



good generalisation

## width effect on landscape

- **narrow network:** chaotic, isolated minima
- **wide network:** smooth, connected valleys

## batch size effect on landscape

- small batch → noisy → finds flat minima
  - large batch → smooth → finds sharp minima
-

# mode connectivity

## the discovery

- **\*observation** different training runs find different minima
- **question** is it possible for to be connected?
  - **answer** YES! (perhaps with slight path elevation)

## types of connectivity

### linear mode connectivity

$$L(\theta_{\text{interpolated}}) \approx L(\theta_1) \approx L(\theta_2)$$

### nonlinear path

$$\theta(t) = \text{bezier\_curve}(\theta_1, \theta_{\text{curve}}, \theta_2, t)$$

low loss path exists but curves

### implication

The solution space is a connected manifold, not isolated points!

---

---

## deep learning success story (dl-ss)

### why deep learning works?

1. **Representation Learning**: learns what to look for
2. **over-parameterisation**: makes optimisation easier
3. **foundation models**: Amortises learning across tasks
4. **modern optimisers**: Adaptive, stable, fast
5. **implicit regularization**: SGD prefers simple solutions
6. **double descent**: More is actually better
7. **good geometry**: Loss landscapes are navigable

### the surprising truth

"We don't train networks in spite of having too many parameters, we train them BECAUSE we have too many parameters!"

# Lottery Ticket Hypothesis

## The Surprising Discovery

### The Observation (Frankle & Carbin, 2019)

"A randomly initialised, dense neural network contains a subnetwork that is initialised such that— when trained in isolation — it can match the test accuracy of the original network after training for at most the same number of iterations."

#### in simple terms

Your neural network is like a lottery with millions of tickets:

- **most tickets lose** (bad subnetworks)
- **a few tickets win** (good subnetworks)
- **training finds the winning tickets**
- **the winning ticket was there from the start!**

#### why this may matter?

- we can remove 90-95% of parameters
- but ONLY if we keep the right initialisation
- suggests structure emerges very early in training
  - structure, i.e. the winning subnetwork

---

## the Lottery Ticket algorithm

### step-by-step

```
def find_winning_ticket(network, data):  
    # Step 1: Random Initialization  
    θ₀ = random_initialize()  
  
    # Step 2: Train to Completion  
    θ_final = train(network, θ₀, data)  
  
    # Step 3: Identify Important Weights (Pruning)  
    # pruning can be e.g. magnitude
```

```

mask = magnitude_prune(theta_final, sparsity=90%)

# Step 4: Reset to Original Init with Mask
theta_ticket = mask ⊙ theta_0 # Element-wise product

return theta_ticket, mask

```

## mathematical formulation:

Given network  $f(x; \theta)$  with parameters  $\theta \in \mathbb{R}^n$ :

1. **initialize**:  $\theta_0 \sim \mathcal{N}(0, \sigma^2)$
2. **train**:  $\theta_T = \text{SGD}(f, \theta_0, \mathcal{D}, T)$
3. **prune**:  $m_i = \mathbb{1}[|\theta_{T,i}| > \tau]$  where  $\tau$  is the pruning threshold
4. **reset**:  $\theta_{\text{ticket}} = m \odot \theta_0$

## key finding

$\mathcal{L}(f(x; m \odot \theta_0)) \approx \mathcal{L}(f(x; \theta_T))$  ✓ Works if initialised to original values !  
 $\mathcal{L}(f(x; m \odot \theta'_0)) \gg \mathcal{L}(f(x; \theta_T))$  ✗ Fails if reinitialised!

---

## why do lottery tickets exist?

### the combinatorial argument

#### number of possible subnetworks

$$N_{\text{subnets}} = \binom{n}{k} \approx \frac{n^k}{k!}$$

where  $n$  = total parameters,  $k$  = parameters kept

#### probability at least one is good

$$P(\text{winning ticket exists}) = 1 - (1 - p)^{N_{\text{subnets}}}$$

## Concrete Example:

- Network with 1M parameters, keeping 10% (100K)
- Number of possible subnetworks:  $\binom{10^6}{10^5} \approx 10^{250,000}$
- Even if  $p = 10^{-100}$  per subnet, we still have winners!

## The Initialization Lottery:

```
Dense Network (1M params)
↓
Contains ~10^250,000 subnetworks
↓
Most are bad, but some are golden!
↓
Training identifies which ones
```

---

## pruning strategies

### magnitude pruning (most common)

remove weights with smallest absolute values:

$$m_i = \begin{cases} 1 & \text{if } |\theta_i| \geq \tau_p \\ 0 & \text{if } |\theta_i| < \tau_p \end{cases}$$

where  $\tau_p$  is the  $p$ -th percentile of  $|\theta|$

### structured vs unstructured

**unstructured** (Original LTH):

- prune individual weights
- maximum flexibility
- hard to accelerate

**structured**:

- prune entire neurons / channels / layers
- less flexible
- hardware-friendly

### iterative magnitude pruning (IMP)

Better for high sparsity:

```
for round in range(num_rounds):
    train_model()
    prune_20_percent() # Gradual
    reset_to_original_init()
```

---

## when do tickets emerge?

### early-bird tickets (You et al., 2020)

**question:** when, during training, does the winning structure emerge?

**finding:** very early, often within 10-20% of training!

### the rewinding insight

instead of resetting to iteration 0, reset to iteration  $k$ :

$$\theta_{\text{ticket}} = m \odot \theta_k \quad \text{where } k \ll T$$

**results:**

- $k = 0$  (original init): Works up to ~90% sparsity
- $k = 0.1 T$  (10% trained): Works up to ~98% sparsity
- $k = 0.5 T$  (half trained): Almost any sparsity works

### implication:

the "lottery ticket" isn't completely random—early training refines it!

---

## lottery tickets at scale

### BERT and Transformers

**challenge:** original LTH often fails for large models!

**solution:** Late rewinding

- Don't reset to iteration 0
- Reset to ~1000 steps in
- Now works for BERT, GPT-2

## the Lottery Ticket Hypothesis 2.0

for large models:

1. **phase 1** (0-1000 steps): Find good basin
2. **phase 2** (1000+ steps): Find specific minimum
3. **pruning**: Identifies important structure
4. **rewinding**: Keep Phase 1, redo Phase 2

## Practical Impact

- 90% sparse BERT matches dense performance
  - 10x inference speedup possible
  - transfer: Tickets found on one task work for others!
- 

## what Lottery Tickets tell us?

1. **initialisation contains solution**
  - networks don't learn from blank state
  - good subnetworks exist at initialisation
2. **parameterisation = more tickets**
  - larger networks → more subnetworks → higher chance of winners
  - may explains why bigger models train better
3. **structure emerges early**
  - important connections identified quickly
  - rest of training is refinement

## Practical Applications:

- network compression**: prune 90%+ parameters
- faster training**: train only the winning ticket
- architecture search**: tickets reveal important structures
- transfer learning**: tickets generalise across tasks

[Lottery Ticket Hypothesis](#)

---

---

## Neural Tangent learning - understanding deep learning through kernel methods

### the core idea

### the revolutionary insight (Jacot et al., 2018)

"In the infinite width limit, neural network training dynamics become linear and are governed by a fixed kernel—the Neural Tangent Kernel"

### what this means:

**Finite Width Network (Normal):**

- weights change during training

- features evolve and adapt
- non-linear dynamics
- hard to analyse

**Infinite Width Network (NTK):**

- features stay (approximately) fixed
- training = linear regression in feature space
- dynamics completely predictable
- becomes a kernel method!

**the bridge:**

Neural Networks  $\longleftrightarrow$  Kernel Methods

---

## mathematical foundation

### The Neural Tangent Kernel Definition

for neural network  $f(x; \theta)$  with parameters  $\theta$ :

$$\Theta(x, x') = \langle \nabla_{\theta} f(x; \theta), \nabla_{\theta} f(x'; \theta) \rangle = \sum_{i=1}^p \frac{\partial f(x)}{\partial \theta_i} \frac{\partial f(x')}{\partial \theta_i}$$

**intuition:** how similarly does the network change at inputs  $x$  and  $x'$ ?

### training dynamics in function space

under gradient flow ( $\dot{\theta} = -\nabla_{\theta} \mathcal{L}$ ):

$$\frac{df(x, t)}{dt} = - \sum_{i=1}^n (f(x_i, t) - y_i) \cdot \Theta(x, x_i)$$

in matrix form:

$$\frac{d\vec{f}}{dt} = -\Theta \cdot (\vec{f} - \vec{y})$$

this is a LINEAR ODE! Solution:

$$f(x, t) = f(x, 0) + \Theta(x, X)^{\top} (\mathbf{I} - e^{-\Theta t}) \Theta^{-1} (\vec{y} - f(X, 0))$$


---

## the infinite width limit

### what happens as width $\rightarrow \infty$

consider a 2-layer network:

$$f(x) = \frac{1}{\sqrt{m}} \sum_{j=1}^m a_j \sigma(w_j^T x)$$

where  $m$  = width,  $a_j$  = output weights,  $w_j$  = hidden weights

### key results:

#### 1. at initialization (width $\rightarrow \infty$ ):

- $f(x) \rightarrow \mathcal{GP}(0, \Sigma)$  (Gaussian Process)
- $\Theta(x, x') \rightarrow \Theta^*(x, x')$  (deterministic kernel)

#### 2. during training (width $\rightarrow \infty$ ):

- $\Theta(x, x', t) \approx \Theta(x, x', 0)$  (kernel stays constant!)
- features don't evolve
- training = kernel regression

## the width scaling:

$$\text{Feature learning} \sim \frac{1}{\sqrt{\text{width}}}$$

$$\text{Kernel stability} \sim 1 - \frac{1}{\sqrt{\text{width}}}$$

---

## NTK for Different Architectures

### Fully Connected Networks

$$\Theta_{\text{FC}}(x, x') = \sum_{l=0}^L \Sigma^{(l)}(x, x') \cdot \dot{\Sigma}^{(l)}(x, x')$$

where  $\Sigma^{(l)}$  is the kernel at layer  $l$

### Convolutional Networks

$$\Theta_{\text{CNN}}(x, x') = \sum_{\text{patches}} \Theta_{\text{FC}}(\text{patch}(x), \text{patch}(x'))$$

CNNs = Locally connected NTK

## Attention/Transformers

$$\Theta_{\text{Attn}}(x, x') = \sum_{i,j} \alpha_{ij}(x)\alpha_{ij}(x') \cdot \Theta_{\text{base}}(x_i, x'_j)$$

where  $\alpha_{ij}$  are attention weights

## Key Pattern:

Architecture → Kernel Structure → Inductive Bias

---

## finite width effects

### real networks are not infinite!

**finite width** (real world):

- NTK evolves during training
- features learn and adapt
- better generalization

**the feature learning regime:**

$$\Theta(t) = \Theta(0) + \underbrace{\Delta\Theta(t)}_{\text{Feature Learning}}$$

$$\Delta\Theta \propto \frac{1}{\sqrt{\text{width}}}$$

## the transition

Narrow Networks (width < 100)	Wide Networks (width ~ 1000)	Infinite Networks (width → ∞)
↓	↓	↓
Strong Feature Learning	Moderate Feature Learning	No Feature Learning
↓	↓	↓
Hard to Train	Sweet Spot	Kernel Regression

## optimal width

not too narrow (hard optimization)  
not too wide (poor feature learning)

---

## why NTK matters

### Theoretical Insights

- ✓ **explains trainability:** wide networks have better-conditioned NTK
- ✓ **predicts convergence:** can compute convergence rate from  $\lambda_{\min}(\Theta)$
- ✓ **architecture design:** NTK reveals inductive bias
- ✓ **generalization bounds:** kernel theory applies

### practical applications

#### 1. infinite-width networks as baselines

```
# If infinite-width does well → architecture is good  
# If finite-width much better → feature learning matters
```

#### 2. architecture search via NTK

- Compute NTK at initialization
- Good condition number → trainable architecture

#### 3. understanding scale

- When is a network "wide enough"?
- NTK condition number tells us

### limitations

- ✗ real networks learn features (good thing!)
  - ✗ NTK can't explain all phenomena (e.g., grokking)
  - ✗ infinite width is impractical
- 

## NTK vs Lottery Tickets

two perspectives on the same phenomenon

Aspect	Lottery Ticket	Neural Tangent Kernel
Focus	Sparse structure	Training dynamics
Key insight	Good subnetworks exist	Wide → linear dynamics
Initialization	Contains solution	Determines kernel

Aspect	Lottery Ticket	Neural Tangent Kernel
Overparameterization	More tickets	Better conditioning
Feature learning	Implicit	Decreases with width
Practical use	Pruning, compression	Architecture design

## the unified view

both explain why over-parameterisation helps:

**Lottery Tickets**: more parameters → more subnetworks → higher chance of winners

**NTK**: more parameters → smoother kernel → easier optimization

## the paradox resolved

- **classical view**: parameterisation → Overfitting X
- **modern view**: parameterisation → Better optimisation → Better generalisation ✓

both LTH and NTK explain different aspects of why this works!

---



---

## Lottery Ticket and NTK summary and key takeaways

### Lottery Ticket Hypothesis

 **core Message**: Your network already contains the solution at initialization

**remember:**

- magnitude pruning after training
- reset to original initialization
- 90%+ sparsity is achievable
- structure emerges early

**practical Impact**: Compression, efficiency, understanding

### Neural Tangent Kernel

 **core Message**: Wide networks train like linear models


**remember:**

- $\Theta(x, x') = \langle \nabla f(x), \nabla f(x') \rangle$
- infinite width → fixed kernel

- finite width → feature learning
- architecture → kernel → bias

**Practical Impact:** Theory, architecture design, understanding

## the big picture



both lead to: **easier training + better generalization!**

---

---

## Some resources used

[Noise, a flaw in human judgement] Daniel Kahneman, Oliver Sibony, Cass R. Sunstein, 2021