


05 Attention and Transformers

Today's Roadmap

1. From RNNs to Attention: The motivation
 2. The Attention Mechanism: Core concepts and mathematics
 3. Transformer Architecture: Building blocks
 4. Why Transformers Work: Theoretical insights
 5. Sequential Data Representation: Beyond NLP
 6. Architecture Deep Dive: Implementation details
 7. Computational Considerations: Efficiency and scaling
-
-

1. The Problem with Recurrent Architectures

The Sequential Bottleneck



RNN/LSTM Processing:

$$\begin{array}{ccccccc} h_1 & \rightarrow & h_2 & \rightarrow & h_3 & \rightarrow & h_4 & \rightarrow \dots \rightarrow h_n \\ \uparrow & & \uparrow & & \uparrow & & \uparrow \\ x_1 & & x_2 & & x_3 & & x_4 & & x_n \end{array}$$

Problems:

- Sequential computation (can't parallelize)
- Long-range dependencies vanish
- Information bottleneck through hidden state

Training RNN

1. training translates to *expansion* of the recursive structure to a multilayer perceptron

2. now we can use gradient descent algorithm

$$L(x_1, \dots, x_T, y_1, \dots, y_T, w_h, w_o) = \frac{1}{T} \sum_{t=1}^T l_t(y_t, o_t)$$

$$\begin{aligned} 3. \quad \frac{\partial L}{\partial w_h} &= \frac{1}{T} \sum_{t=1}^T \frac{\partial l(y_t, o_t)}{\partial w_h} = \frac{1}{T} \sum_{t=1}^T \frac{\partial l(y_t)}{\partial o_t} \frac{\partial g(h_t, w_o)}{\partial h_t} \frac{\partial h_t}{\partial w_h} \\ &\frac{\partial h_t}{\partial w_h} = \frac{\partial f(x_t, h_{t-1}, w_h)}{\partial w_h} + \frac{\partial f(x_t, h_{t-1}, w_h)}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial w_h} \end{aligned}$$

4. last $\frac{\partial h_t}{\partial w_h}$ (last part) on the weights from the last state h_{t-1} and, at the same time w_h
 (first part after equation)
5. it is possible to make some approximations, but it is still not a stable computation
6. it needs to be computed only on some last T states only

The Vanishing Gradient Problem Revisited

Even with LSTM/GRU:


- Path length between distant tokens grows linearly: $O(n)$
- Gradient flow diminishes over long sequences
- Earlier tokens have exponentially smaller influence


Mathematical View:

$$\frac{\partial h_t}{\partial h_0} = \prod_{i=1}^t \frac{\partial h_i}{\partial h_{i-1}}$$

If $\frac{\partial h_i}{\partial h_{i-1}} < 1$, gradients vanish as $t \rightarrow \infty$

- partially solved in recurrent models like **GRU** and **LSTM**
 - introduction of ReLU activation
 - gating





- RNN
 - sequential flow
 - diminishing gradient
 - long path with no shortcuts
- Transformer
 - flow by layers with residual connections
 - stable gradients with slight degradation

The Fixed Context Problem

Encoder-Decoder RNN:

```

# Entire input compressed into single vector c
encoder_output = encode(x1, x2, ..., xn) # → c (fixed size!)
decoder_output = decode(c, y1, y2, ..., ym)

# Problem: c must capture EVERYTHING about the input
# This is an information bottleneck!

```

“rnn-encoder-decoder.svg” could not be found.

Example: Translation

English: "The agreement on the European Economic Area was signed in August 1992"

[14 words → compressed to fixed vector c → decode]

For each output word, decoder sees:

- Same context vector c
- No direct access to specific input words
- Can't "look back" at relevant parts

2. Attention

"Instead of encoding the entire input into a fixed context vector, let the decoder **attend** to different parts of the input at each decoding step."

The Problem: Fixed Context Bottleneck

- before attention
 - encoder-decoder compressed the **entire input sequence into a single fixed-size context vector****
 - e.g. "*The agreement on the European Economic Area was signed in August 1992*"
 - all 14 words had to be squeezed into one vector
 - severe information bottleneck, especially for long sequences
 - losing important relations between words

The Bahdanau Innovation

[Bahdanau et al. \(2014\)](#) (over 40 thousand citations) first attention mechanism for neural machine translation.

- instead of using a fixed context vector,
- decoder **dynamically attends to different parts of the encoder output** at each decoding step:

1. **Encoder produces a sequence of hidden states:** h_1, h_2, \dots, h_n (one per input word)
2. **At each decoder step t , compute alignment scores** between decoder state s_{t-1} and each encoder state

$$h_j : e_{tj} = \text{score}(s_{t-1}, h_j) = v^T \tanh(W_1 s_t + W_2 h_j)$$

3. Convert to attention weights via softmax:


$$\alpha_{tj} = \frac{\exp(e_{tj})}{\sum_{k=1}^n \exp(e_{tk})}$$

4. Compute context vector as weighted sum:

$$c_t = \sum_{j=1}^n \alpha_{tj} h_j$$

5. **Use c_t for decoding** : Generate output word based s_t on both decoder state and context

6. It is possible to build two series of hidden states: h_1, h_2, \dots, h_n (from first to last word) and h_n, h_{n-1}, \dots, h_1 (from last to first) and combine them



Example in translation:

Translating: "The cat sat" → "Le chat s'est assis"

When generating "chat" (French for cat):

- Decoder computes attention over ["The", "cat", "sat"]
- High attention weight on "cat": $\alpha = [0.1, 0.8, 0.1]$

- Context vector emphasizes "cat" representation
- Decoder uses this to generate correct translation

Model training

1. specify the encoder and decoder models (LSTM, GRU)
2. add attention module
3. build a translation module (sequence-to-sequence model too)
4. train all

Why It Mattered

Impact on the field:

- **Solved the bottleneck**: No need to compress entire sequence into single vector
- **Enabled long sequences**: Could handle 50+ word sentences effectively
- **Interpretable**: Attention weights showed which source words influenced each target word
- **State-of-the-art results**: Dramatically improved translation quality

Difference from Transformers:

- **Still used RNNs**: Attention was an *add-on* to RNN encoder-decoder
- **Sequential processing**: Had to process words one at a time (no parallelization)
- **Auxiliary mechanism**: Attention helped RNNs, but RNNs were still the core architecture

The path to Transformers:

Bahdanau Attention (2014): "Let's add attention TO RNNs"

↓

Huge improvement!

↓

Vaswani et al. (2017): "What if we use ONLY attention?"

↓

Transformer is born

Mathematical Formulation

The Bahdanau attention mechanism can be written as:

$$context_t = \sum_{j=1}^n softmax(score(s_t, h_j)) \cdot h_j$$

where the score function learns to measure

alignment between decoder state s_t and encoder state h_j

This is conceptually similar to the Transformer's attention:

$$\text{Attention}(Q, K, V) = \text{softmax}(QK^T)V$$

but

- Bahdanau: Single query (decoder state) attending to sequence
 - Transformer: All positions attend to all positions simultaneously
 - Bahdanau: Score computed via learned MLP
 - Transformer: Score computed via dot product (simpler, faster)
-

Bahdanau attention

- proved that explicit attention mechanisms could dramatically improve sequence-to-sequence models.
 - showed that models could "look back" at relevant input positions rather than relying on a compressed context
 - inspired the Transformer, which took the idea further by making attention the *only* mechanism, removing RNNs entirely.
-
-

Transformer explained

Attention Intuition

Translation: "The cat sat on the mat" → "Le chat s'est assis sur le tapis"

When generating "chat":

Attention weights: [0.8, 0.15, 0.02, 0.01, 0.01, 0.01]
 ↓ ↓ ↓ ↓ ↓ ↓
 chat s'est assis sur le tapis

Model focuses on "cat" (0.8 weight) while translating!

From Implicit to Explicit Attention

- **RNN Encoder-Decoder (2014):**
 - Implicit attention through recurrent connections
 - Each h_t depends on all previous
- $$h_t = f(h_{t-1}, x_t)$$

- Bahdanau Attention (2015):

- Explicit attention weights

$$\alpha_t = \text{softmax}(\text{score}(h_t, \text{encoder-states}))$$

$$\text{context}_t = \sum \alpha_t[i] * \text{encoder-states}[i]$$

- Self-Attention / Transformer (2017):


- Attention as the PRIMARY mechanism (not auxiliary)
- No RNN at all

$$\text{output} = \text{Attention}(Q, K, V)$$

5. Attention Mechanism

- model needs to cope with input passages (e.g. in NLP) of different lengths
- know connections between elements that depend on attention
- transformer gets both by using **dot-product self-attention**

Self-Attention: The Core Operation



(from Understanding deep learning)

- standard network computes a linear transformation with a non-linear function
- **self-attention $sa[]$** takes
 - N inputs
 - returns N output vectors of the same size (e.g., in NLP a word or some sub-word)
- a **value** is computed from each token


$$v_m = \beta_v + \Omega_v x_m$$

using biases β and weights Ω

- may be done in parallel
- a scalar $a[x_i, x_j]$ is the **attention** that token x_j pays to token x_i
 - attentions $a[\cdot, x_j]$ sum up to 1
- **self attention**

$$sa_j(x_1, \dots, x_N) = \sum_{i=1}^N a[x_i, x_j] \cdot v_i$$

- a weighted sum of all values v_i
- weights $a(\cdot, x_n)$ are non-negative and sum-up to 1
- each self-attention $sa_i[x_1, \dots]$ can be thought as routing of the original N tokens with different proportions for the current task
- all can be computed in parallel per token
- in the figure above
 - N inputs are taken
 - in the left-most a routing to $sa_1(x)$ is computed with weights 0.1, 0.3, 0.6
 - then two different **routings**



(from Prince, Understanding deep learning, MIT, 2023)

- *query* vectors are computed as $q_n = \beta_q + \Omega_q x_n$

- key vectors are computed as $k_n = \beta_k + \Omega_q x_n$
- dot products are passed to a softmax giving attention values

Given input $\mathbf{X} \in \mathbb{R}^{N \times D}$:


$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Where:

- $Q = XW_Q$ (Queries): "What am I looking for?"
- $K = XW_K$ (Keys): "What information do I have?"
- $V = XW_V$ (Values): "What information to aggregate?"
where dimensions are
 - d model dimension, size of the embedding,
 - d_k the key and query dimension: $d_k = d/n_{heads}$ (standard relationship)
 - d_v dimension of values in attention: $d_v = d/n_{heads}$ (typically $d_v = d_k$)
 - typically embedding dimension needs to be a multiple of the number of heads
 - if $d_v \neq d_k$, an additional attention may be used where additional trained matrices map both to the same dimension
- $W_Q \in \mathbb{R}^{d \times d_v}$, $W_K \in \mathbb{R}^{d \times d_k}$, $W_V \in \mathbb{R}^{d \times d_v}$
 - dot products in computation may get to large values
 - and the softmax regions where largest value dominates, and the gradients get very small
 - model gets hard to train
 - the scaling by the square root of the number of rows of keys (and queries) prevents it


Key, query, value

- the **dot product** returns a measure of similarity between its arguments
- weights $a[x., x_n]$ depend on relative similarities between the n-th query and all the keys
- queries and keys should have the same dimensions
 - * there is a possibility of mapping to a common value



Multiple attention heads

- multiple heads may be computed in parallel



- typically for model dimension d and num_{heads} heads, the values, queries and keys will **all be of the same size** allowing for efficient computation
- multiple heads are concatenated

3. Understanding Query, Key, and Value: The Heart of Attention

Beyond the Matrices: Building Intuition

[Transformer explainer](#)

The Fundamental Question

Why do we need THREE separate matrices (Q, K, V)?

Bad answer: "Because the math works out"

- in my (/igor) opinion the [Attention Is All You Need, Vaswani et al., 2017](#)) paper caused Transformers a great deal of harm
- people now understand Transformer just as clever multiplication of matrices
- almost nobody asks themselves how the model "thinks", stores knowledge

Good answer: "Because attention is answering three distinct questions:"

1. **Query (Q):** "What information am I looking for?"
 2. **Key (K):** "What information do I offer?"
 3. **Value (V):** "What is the actual information I provide?"
-

Library Search

You walk into a library looking for:
QUERY: "Books about neural networks written after 2017"

The librarian's process:

1. **Check catalog cards (KEYS):**
 - Book 1: "Deep Learning, 2016, Goodfellow"
 - Book 2: "Neural Networks, 2020, Smith"
 - Book 3: "Transformers, 2017, Vaswani"
2. **Match your query against keys:**
 - Book 1: (too old)
 - Book 2: (perfect match!)
 - Book 3: (borderline - exactly 2017)
3. **Retrieve actual books (VALUES):**
 - The catalog card (KEY) tells you WHERE to look
 - But what you get is the actual BOOK (VALUE)
 - Key = "metadata for matching"
 - Value = "the content you actually want"

In attention:

```
Your query: "Books about neural networks after 2017"
↓
Similarity scores: [0.1, 0.9, 0.5] # How well each book matches
↓
Weighted retrieval: 0.1 * Book1 + 0.9 * Book2 + 0.5 * Book3
```

↓

You get: Mostly Book 2, some Book 3, almost no Book 1

[FIGURE NEEDED: Visual diagram of library search with Q, K, V highlighted]

Database Query (SQL-like)

```
SELECT value_column
FROM table
WHERE key_column MATCHES query
ORDER BY similarity(key_column, query) DESC

-- In attention, this becomes:
-- Weighted retrieval based on soft matching!
```

Example: Student Database

```
# You want information about: "CS students with GPA > 3.5"
query = [field: "CS", criteria: "GPA > 3.5"]

# Database entries:
students = [
    {key: "Alice, CS, 3.9",      value: "Research: NLP, Skills: PyTorch"}, 
    {key: "Bob, Math, 3.8",      value: "Research: Topology, Skills: Proofs"}, 
    {key: "Carol, CS, 3.2",      value: "Research: Graphics, Skills: OpenGL"}, 
    {key: "Dave, CS, 3.7",       value: "Research: ML, Skills: TensorFlow"}]

# Attention scores (how well keys match query):
scores = [0.95, 0.1, 0.3, 0.9] # Alice and Dave match best

# Weighted retrieval (soft SQL):
result = 0.95 * "Research: NLP..." + 0.1 * "Research: Topology..." +
        + 0.3 * "Research: Graphics..." + 0.9 * "Research: ML..."

# You get: Mixture of Alice and Dave's info (CS students with high GPA)
```

Key insight:

- **Keys** are for matching/indexing
- **Values** are the actual content you retrieve
- **Query** is what you're searching for

Search Engine

User types: "best restaurants near me" ← QUERY

↓

Google's index:

Document 1:

KEY: [location:nearby, topic:restaurants,
rating:4.5, recency:2024]

VALUE: "Mario's Pizza: 123 Main St,
4.5★, Italian cuisine..."

Document 2:

KEY: [location:nearby, topic:hotels,
rating:4.0, recency:2023]

VALUE: "Grand Hotel: 456 Oak Ave,
4.0★, Luxury accommodation..."

Matching process:

1. **Encode query:** "restaurants nearby" → query vector
2. **Compare to keys:** Which documents are relevant?
 - Doc 1: topic=restaurants ✓, location=nearby ✓ → HIGH score
 - Doc 2: topic=hotels ✗, location=nearby ✓ → LOW score
3. **Retrieve values:** Return actual document content (weighted by scores)

Why separate Key and Value?

- **Key:** Optimised for fast similarity search
 - Compressed representation: [location, topic, rating, recency]
 - Like a hash/index for quick lookup
- **Value:** The full, rich information
 - Complete document content
 - Much larger, more detailed
 - You don't search through this directly (too expensive!)

The Mathematical Picture: Information Retrieval

From Retrieval to Continuous Attention

Hard retrieval (traditional):

```
def hard_lookup(query, keys, values):
    # Find BEST matching key
    best_idx = argmax(similarity(query, keys))
    # Return corresponding value
    return values[best_idx]
```

Problem: Discrete, non-differentiable

Soft retrieval (attention):

```
def soft_lookup(query, keys, values):
    # Compute similarity to ALL keys
    scores = similarity(query, keys) # [s1, s2, ..., sn]

    # Convert to probabilities (softmax)
    weights = softmax(scores) # [w1, w2, ..., wn]

    # Weighted average of ALL values
    return Σ wi * values[i]

# Differentiable! Can learn what to attend to
```

- most Transformer architectures use some form of top-k architecture
 - in NN top-k means applying the weights modification **only** to the highest **k** outputs
 - thus only the meaningful classes are modified
 - it can be done using differentiable operations, not just cutting off

Visualisation of scores → weights:

```
Raw scores:      [2.1,      0.3,      -0.5,      1.8]
                  ↓ softmax
Attention weights: [0.58,      0.10,      0.04,      0.28] # Sum to 1.0
                  ↓
Weighted sum:    0.58*V1 + 0.10*V2 + 0.04*V3 + 0.28*V4
```

Translating "The cat sat on the mat" → French

Setup: Decoder generating "chat" (French for "cat")

Semantic dimensions:

1. Is-Animal (0-1)

2. **Is-Action** (0-1)
3. **Is-Object** (0-1)
4. **Definiteness** (0-1)

Step 1: The Query (from decoder)

Decoder's current state: "I'm trying to generate the French word for 'cat'"

```
Q = [0.9, 0.1, 0.2, 0.3]
      |     |     |
      |     |     |     └── definiteness (the cat, not just any cat)
      |     |     |     └── is-object (0.2 - can be object but not here)
      |     |     |     └── is-action (0.1 - not an action)
      |     |     └── is-animal (0.9 - STRONGLY looking for animal!)

# Query says: "I need something that is primarily an ANIMAL,
#                 not an action, possibly definite"
```

Step 2: The Keys (from encoder - what each word "advertises")

Each English word broadcasts what it contains:

```
# "The" (first occurrence)
K_the1 = [0.0, 0.0, 0.0, 1.0]
          |     |     |
          |     |     |     └── definiteness=1.0 (it's "the"!)
          |     |     |
          |     |     └── is-animal=0.0 (not an animal)

# "cat"
K_cat = [1.0, 0.0, 0.3, 0.0]
          |     |     |
          |     |     |     └── definiteness=0.0 (not a determiner)
          |     |     |
          |     |     └── is-object=0.3 (can be object)
          |     |
          |     └── is-action=0.0 (not an action)
          |
          └── is-animal=1.0 (YES! I'm an animal!)

# "sat"
K_sat = [0.0, 1.0, 0.0, 0.0]
          |     |
          |     └── is-action=1.0 (I'm a verb/action!)
          |
          └── is-animal=0.0 (not an animal)

# "on"
K_on = [0.0, 0.0, 0.0, 0.0]
          |
          └── (preposition - low on all semantic features)

# "the" (second occurrence)
K_the2 = [0.0, 0.0, 0.0, 1.0]
```

```

    └ definiteness=1.0

# "mat"
K_mat = [0.0, 0.0, 1.0, 0.0]
    |   |
    |   └ is-object=1.0 (I'm an object!)
    |   └ is-action=0.0
    └ is-animal=0.0 (not an animal)

```

Step 3: Compute Attention Scores ($Q \cdot K$)

$Q = [\text{is-animal}, \text{is-action}, \text{is-object}, \text{definiteness}]$

Element-wise multiplication and sum:

```

# Q · K_the1
score_the1 = 0.9×0.0 + 0.1×0.0 + 0.2×0.0 + 0.3×1.0 = 0.3
                └ no animal match   └ some definiteness

# Q · K_cat
score_cat = 0.9×1.0 + 0.1×0.0 + 0.2×0.3 + 0.3×0.0 = 0.96
                └ STRONG animal match └ minor object match

# Q · K_sat
score_sat = 0.9×0.0 + 0.1×1.0 + 0.2×0.0 + 0.3×0.0 = 0.1
                └ action doesn't match our query

# Q · K_on
score_on = 0.9×0.0 + 0.1×0.0 + 0.2×0.0 + 0.3×0.0 = 0.0
                └ preposition, no
match

# Q · K_the2
score_the2 = 0.9×0.0 + 0.1×0.0 + 0.2×0.0 + 0.3×1.0 = 0.3
                └ same as first
"the"

# Q · K_mat
score_mat = 0.9×0.0 + 0.1×0.0 + 0.2×1.0 + 0.3×0.0 = 0.2
                └ some object match

```

Raw scores: [0.30, 0.96, 0.10, 0.00, 0.30, 0.20]

	the ₁	cat	sat	on	the ₂	mat
	[0.30]	0.96	0.10	0.00	0.30	0.20]

Step 4: Scale by $\sqrt{d_k}$

```
d_k = 4 # dimension of keys; needed for scalability
scaling_factor =  $\sqrt{4} = 2.0$ 

scaled_scores = [0.30, 0.96, 0.10, 0.00, 0.30, 0.20] / 2.0
                = [0.15, 0.48, 0.05, 0.00, 0.15, 0.10]
```

Step 5: Apply Softmax → Attention Weights

```
# exp(scaled_scores)
exp_scores = [exp(0.15), exp(0.48), exp(0.05), exp(0.00), exp(0.15),
exp(0.10)]
              = [      1.16,      1.62,      1.05,      1.00,      1.16,
1.11]

# Normalize (sum = 7.10)
attention_weights = [1.16, 1.62, 1.05, 1.00, 1.16, 1.11] / 7.10
                     = [0.16, 0.23, 0.15, 0.14, 0.16, 0.16]
                     α ≈ [0.16, 0.23, 0.15, 0.14, 0.16, 0.16]
                           the1   CAT     sat     on     the2   mat
```

Visualization:

Word:	the ₁	cat	sat	on	the ₂	mat
Weight:	16%	23%	15%	14%	16%	16%
	—	—	—	—	—	—

↑
Highest attention to "cat"!

Step 6: The Values (actual semantic content to retrieve)

Values contain rich, contextualized information:

```
# V_the1: Determiner introducing "cat"
V_the1 = [0.1, 0.0, 0.0, 0.8]
            |                               ↳ definite article marker
            |                               ↳ minimal semantic content

# V_cat: Rich animal semantics
```

```

V_cat = [0.9, 0.0, 0.1, 0.7]
        |           |           |   └ definite (preceded by "the")
        |           |           |   └ minor object role
        |           |           └ not action
        └ ANIMAL (feline, domestic, pet)

# V_sat: Action/state information
V_sat = [0.0, 0.9, 0.0, 0.3]
        |           └ PAST ACTION (sitting)
        └ not animal

# V_on: Prepositional relationship
V_on = [0.0, 0.0, 0.5, 0.0]
        |           └ spatial relationship
        └ not animal

# V_the2: Determiner for "mat"
V_the2 = [0.0, 0.0, 0.0, 0.9]
        └ definite article

# V_mat: Object semantics
V_mat = [0.0, 0.0, 0.9, 0.6]
        |           └ definite object
        |           └ OBJECT (floor covering)
        └ not animal

```

Step 7: Weighted Sum of Values (Final Output)

```

output = α₁·V_the₁ + α₂·V_cat + α₃·V_sat + α₄·V_on + α₅·V_the₂ + α₆·V_mat
       = 0.16·[0.1, 0.0, 0.0, 0.8] + 0.23·[0.9, 0.0, 0.1, 0.7]
       + 0.15·[0.0, 0.9, 0.0, 0.3] + 0.14·[0.0, 0.0, 0.5, 0.0]
       + 0.16·[0.0, 0.0, 0.0, 0.9] + 0.16·[0.0, 0.0, 0.9, 0.6]

# Dimension 1 (is-animal):
= 0.16×0.1 + 0.23×0.9 + 0.15×0.0 + 0.14×0.0 + 0.16×0.0 + 0.16×0.0
= 0.016 + 0.207 + 0 + 0 + 0 + 0
= 0.223

# Dimension 2 (is-action):
= 0.16×0.0 + 0.23×0.0 + 0.15×0.9 + 0.14×0.0 + 0.16×0.0 + 0.16×0.0
= 0 + 0 + 0.135 + 0 + 0 + 0
= 0.135

# Dimension 3 (is-object):
= 0.16×0.0 + 0.23×0.1 + 0.15×0.0 + 0.14×0.5 + 0.16×0.0 + 0.16×0.9

```

```

= 0 + 0.023 + 0 + 0.070 + 0 + 0.144
= 0.237

# Dimension 4 (definiteness):
= 0.16×0.8 + 0.23×0.7 + 0.15×0.3 + 0.14×0.0 + 0.16×0.9 + 0.16×0.6
= 0.128 + 0.161 + 0.045 + 0 + 0.144 + 0.096
= 0.574

output = [0.223,      0.135,      0.237,      0.574]
          |           |           |           |
          |           |           |           └ Strong definiteness signal
          |           |           |           └ Some object context
          |           |           |           └ Some action context
          |           |           |           └ Animal signal (from "cat")

```

Interpretation of the Output

The enriched representation [0.22, 0.14, 0.24, 0.57] means:

Original query was looking for: ANIMAL

Got back:

- 22% animal features ← from "cat" (23%)
- 14% action features ← from "sat" (15%)
- 24% object features ← from "mat" (16%)
- 57% definiteness ← from "the"s

Decoder now knows:

"Generate an animal word (cat → chat),
it's definite (le/la),
with some action/object context"

Why this helps translation:

- Primary signal: ANIMAL (0.22) → Generate "chat" (cat)
- Definiteness (0.57) → Use "le chat" not just "chat"
- Context from action/object → Past tense, spatial relationship

Visual Summary: The Complete Flow

INPUT: "The cat sat on the mat"

↓

ENCODER REPRESENTATIONS:

Word	Key (for matching)	Value (content)
the ₁	[0.0, 0.0, 0.0, 1.0]	[0.1, 0.0, 0.0, 0.8]
cat	[1.0, 0.0, 0.3, 0.0]	[0.9, 0.0, 0.1, 0.7]
sat	[0.0, 1.0, 0.0, 0.0]	[0.0, 0.9, 0.0, 0.3]
on	[0.0, 0.0, 0.0, 0.0]	[0.0, 0.0, 0.5, 0.0]
the ₂	[0.0, 0.0, 0.0, 1.0]	[0.0, 0.0, 0.0, 0.9]
mat	[0.0, 0.0, 1.0, 0.0]	[0.0, 0.0, 0.9, 0.6]

↓

DECODER QUERY: [0.9, 0.1, 0.2, 0.3]

"Looking for an animal"

↓

ATTENTION SCORES (Q·K):

[0.30, 0.96, 0.10, 0.00, 0.30, 0.20]
↑↑↑ Highest match!

↓

ATTENTION WEIGHTS (softmax):

[0.16, 0.23, 0.15, 0.14, 0.16, 0.16]
↑↑↑ Focus on "cat"

↓

OUTPUT (weighted sum of Values): _____

[0.22, 0.14, 0.24, 0.57]
└ Mostly animal features, definite, some context

↓

DECODER GENERATES: "chat"

1. Keys are for matching

- $K_{\text{cat}} = [1.0, 0.0, 0.3, 0.0]$ says "I'm an animal"
- Query $Q = [0.9, 0.1, 0.2, 0.3]$ matches strongly (score = 0.96)
- Other keys don't match as well

2. Values are for content

- Even though we matched on "cat" via Key
- We retrieve rich Value: [0.9, 0.0, 0.1, 0.7]
- Contains more than just "is-animal" flag
- Has definiteness, context, semantic richness

3. Soft retrieval averages context

- Not just "cat" (23% weight)

- Also gets definiteness from "the" (16% + 16%)
- Some action context from "sat" (15%)
- Some object context from "mat" (16%)

4. Different queries would give different results

- If query was [0.1, 0.9, 0.1, 0.1] (looking for action)
 - Would attend to "sat" instead
 - Would retrieve action semantics
 - Would generate French verb

```
import numpy as np
x = np.array([1, 2, 3])
print(x * 2)
```

Try It Yourself: Different Queries

Exercise: What would happen with this query?

```
Q_action = [0.1, 0.9, 0.1, 0.2]
            |         └ Looking for ACTION (verb)
            |         └ Not looking for animal

# Compute:
score_cat = 0.1×1.0 + 0.9×0.0 + 0.1×0.3 + 0.2×0.0 = ?
score_sat = 0.1×0.0 + 0.9×1.0 + 0.1×0.0 + 0.2×0.0 = ?

# Which word would get highest attention?
# What would the output vector emphasize?
```

[Check changing query weights](#)

Attention weights change

[→ Live Code Demo] ← You click this during lecture

Why Not Just Use One Matrix?

Thought Experiment: What if $Q = K = V$?

```
# Self-attention with  $Q = K = V = X$ 
attention = softmax(X @ X.T) @ X

# This means:
# - Matching criterion = Content itself
# - Retrieved content = Same as matching criterion
```

Problem 1: Conflation of "what to match" and "what to retrieve"

Example from translation:

Word "bank" in: "I went to the bank by the river"

If Key = Value:

- Key says: "I'm the word 'bank' (financial OR river-side)"
- Value says: Same information

But we want:

- Key to say: "I'm a noun that relates to 'river'"
→ Helps matching in this context
- Value to provide: Full semantic content for "river bank"
→ Different from just the word form

Problem 2: Reduced expressiveness

```
# With  $Q = K = V$  (1 matrix):
Parameters:  $d^2$ 

# With  $Q, K, V$  separate (3 matrices):
Parameters:  $3d^2$ 

# More parameters = more expressive power
# Can learn richer representations
```

Problem 3: Asymmetry of roles

Query: "What do I need?" (generated by current state)
Key: "What can I offer?" (advertised by candidates)
Value: "Here's what I actually give" (delivered content)

These are fundamentally different questions!
Forcing them to be the same limits what the model can learn.

The Key-Value Separation: A Deeper Dive

Why is Key ≠ Value Important?

Intuition: Index vs Content

Think of a book:

- **Key** = Chapter titles, keywords, page numbers (the INDEX)
 - Optimized for: Quick scanning, pattern matching
 - Characteristics: Compressed, abstract, matchable
- **Value** = Actual chapter content (the TEXT)
 - Optimized for: Rich information, semantic content
 - Characteristics: Detailed, specific, informative

```
# Book: "Deep Learning"
key_chapter3 = "Chapter 3: Linear Algebra Review"
    # Concise, tells you what's inside

value_chapter3 = """
    Linear algebra is the study of vectors, matrices...
    [20 pages of detailed content]
"""
    # Rich, detailed information
```

In neural networks:


```
# Encoding a sentence: "The black cat"

# Keys: Compressed, abstract representations for matching
K_the  = embed("determiner, definite")
K_black = embed("adjective, color, descriptor")
K_cat   = embed("noun, animal, subject")

# Values: Rich semantic embeddings with full context
V_the  = embed("article with context from 'black cat'")
V_black = embed("color=black, modifies=cat, visual_property")
V_cat   = embed("animal=feline, color=black, subject=yes, ...")
    # Much richer! Includes compositional information
```

Key vs Value: Index vs Content

Why we separate what we match on (Key) from what we retrieve (Value)



Analogy:

```
# Traditional dictionary
dictionary = {
    "cat": "a small domesticated carnivorous mammal",
    "dog": "a domesticated carnivorous mammal",
    ...
}
result = dictionary["cat"] # Hard lookup: exact match

# Attention as soft dictionary
queries = ["ct", "catt", "dog"] # Fuzzy queries
keys = ["cat", "dog", "bird"]
values = [embedding_cat, embedding_dog, embedding_bird]

# Each query attends to ALL keys, weighted by similarity
result = Σ similarity(query, key_i) * value_i
```

Step-by-Step Computation

Given input sequence $X = [x_1, x_2, \dots, x_n] \in \mathbb{R}^{n \times d}$:

Step 1: Linear Projections

$$Q = XW^Q, \quad K = XW^K, \quad V = XW^V$$

where $W^Q, W^K \in \mathbb{R}^{d \times d_k}$ and $W^V \in \mathbb{R}^{d \times d_v}$

Step 2: Compute Attention Scores

$$\text{scores} = QK^T \in \mathbb{R}^{n \times n}$$

Element (i, j) measures how much query i should attend to key j

Step 3: Scale (Important!)

$$\text{scaled scores} = \frac{QK^T}{\sqrt{d_k}}$$

Why divide by $\sqrt{d_k}$?

$$q \cdot k = \sum_i^d q_i k_i$$

- q_i and k_i have 0 mean and 1 variance
 - $q_i k_i$ has variance 1 too
 - the sum's variance is d and standard deviation \sqrt{d}
 - if scaled $q \cdot k / \sqrt{d}$, the whole variance is 1 too regardless of dimension
 - division keeps gradients stable (prevents softmax saturation)
-

Step 4: Apply Softmax

$$\alpha = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right)$$

Each row sums to 1: $\sum_j \alpha_{ij} = 1$

Step 5: Weighted Sum of Values

$$\text{Output} = \alpha V$$

Attention Score Interpretation

$$\text{score}_{ij} = q_i \cdot k_j = |q_i| |k_j| \cos(\theta_{ij})$$

- High score → query i and key j are "aligned"
- Low score → query i doesn't need key j

After softmax:

$$\alpha_{ij} = \frac{\exp(\text{score}_{ij})}{\sum_k \exp(\text{score}_{ik})}$$

Attention weights form a probability distribution over source positions

Self-Attention vs Cross-Attention

Type	Q from	K,V from	Use Case
Self-Attention	Same sequence	Same sequence	Encoding contextual relationships
Cross-Attention	Target sequence	Source sequence	Encoder-Decoder connection

Self-Attention:

```
# All positions attend to all other positions in same sequence
Q = K = V = X # Same source
# Learns internal structure and dependencies
```

Cross-Attention:

```
# Decoder attends to encoder
Q = decoder_states # What decoder is looking for
K = V = encoder_states # Information from encoder
# Connects source and target sequences
```

5. Multi-Head Attention

The Motivation

Single attention head:

- Learns one notion of "relatedness"
- May miss different types of relationships

Example:

Sentence: "The bank by the river is steep"

Head 1 might learn: syntactic relationships (subject–verb)
Head 2 might learn: semantic relationships (bank–river)
Head 3 might learn: positional relationships (nearby words)

Multi-Head Attention Formula

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$

where each head:

$$\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$

where

- h : number of heads (typically 8 or 16)
 - $W_i^Q, W_i^K \in \mathbb{R}^{d_{model} \times d_k}$ where $d_k = d_{model}/h$
 - $W_i^V \in \mathbb{R}^{d_{model} \times d_v}$ where $d_v = d_{model}/h$
 - $W^O \in \mathbb{R}^{hd_v \times d_{model}}$
-

Why Multiple Heads?

Mathematical Intuition:

Instead of one d_{model} -dimensional attention:

$$\text{Attention}(Q, K, V) \in \mathbb{R}^{n \times d_{model}}$$

Split into h heads of dimension $d_k = d_{model}/h$:

$$\text{head}_i \in \mathbb{R}^{n \times d_k}$$

1. **Subspace specialization**: Each head can attend to different aspects
2. **Parameter efficiency**: h small projections vs 1 large projection
3. **Ensemble effect**: Multiple attention patterns averaged

Different attention patterns across heads [Transformer explained](#)

Computational View

```
# Pseudocode for Multi-Head Attention

def multi_head_attention(Q, K, V, num_heads=8):
    d_k = d_model // num_heads

    # Split into heads
    Q_heads = split_heads(Q, num_heads)  # (batch, heads, seq, d_k)
    K_heads = split_heads(K, num_heads)
    V_heads = split_heads(V, num_heads)

    # Parallel attention for each head
    attention_outputs = []
    for i in range(num_heads):
        head_i = scaled_dot_product_attention(
            Q_heads[:, i], K_heads[:, i], V_heads[:, i]
        )
        attention_outputs.append(head_i)

    return attention_outputs
```

```

# Concatenate and project
concat = concatenate(attention_outputs) # (batch, seq, d_model)
output = linear(concat, W_0)

return output

```

[Multihead attention in PyTorch](#)

Additional attention

If queries and keys do not have the same dimension, we may use an **additional** attention

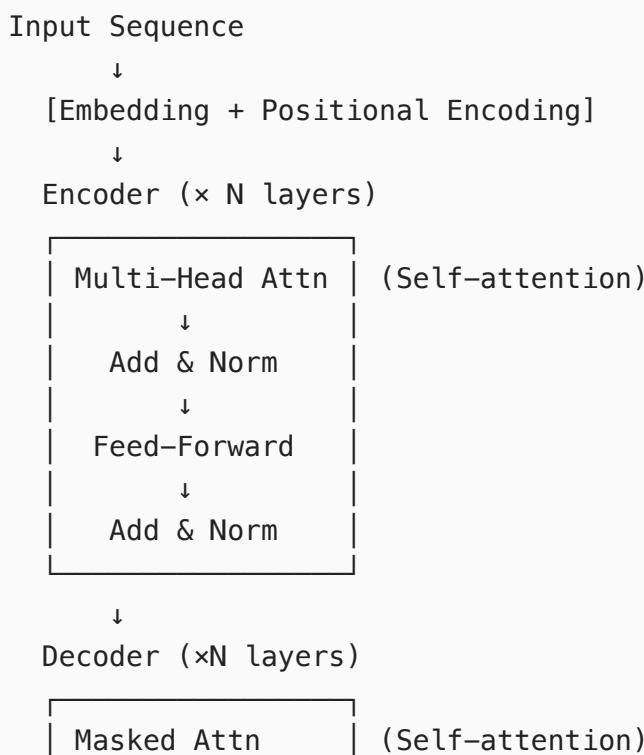
$$attn(q, k) = w_v \tanh(W_q q + W_k k)$$


6. The Transformer Architecture


Three types of transformer architecture

- **encoder** transforms the embeddings into some representation that would support some processing task
- **decoder** predicts the next token to continue the input
- **encoder-decoder** for conversion of one sequence into another e.g., translation

High-Level View








(after Vaswani et al.)

Full Transformer encoder (left) and decoder (right) architecture (one layer both)

- masked multi-head attention masks future tokens preventing glimpses "to the future"



Transformer Encoder Layer

$$\text{EncoderLayer}(X) = \text{FFN}(\text{LayerNorm}(X + \text{MultiHead}(X, X, X)))$$

Detailed Steps:

1. **Self-Attention with Residual:**

$$Z = X + \text{MultiHead}(X, X, X)$$

2. **Layer Normalization:**

$$Z' = \text{LayerNorm}(Z)$$

3. **Feed-Forward with Residual:**

$$H = Z' + \text{FFN}(Z')$$

4. **Layer Normalisation:**

$$\text{Output} = \text{LayerNorm}(H)$$

Feed-Forward Network (FFN)

Architecture:

$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$

Properties:

- Applied **position-wise** (same FFN for each position independently)
- Two linear transformations with ReLU activation
- Typical dimensions: $d_{model} = 512$, $d_{ff} = 2048$

Why FFN after Attention?

1. **Non-linearity:** Attention is linear operations + softmax
 2. **Mixing information:** Attention aggregates, FFN processes
 3. **Capacity:** Adds parameters for complex transformations
-

FFN as 1×1 Convolution

The Position-Wise Feed-Forward Network

Standard FFN Definition

$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$

where:

- $x \in \mathbb{R}^{d_{model}}$ is a single position's representation
- $W_1 \in \mathbb{R}^{d_{model} \times d_{ff}}$, typically $d_{ff} = 4 \cdot d_{model}$
- $W_2 \in \mathbb{R}^{d_{ff} \times d_{model}}$ squeezes back to d_{model}

Key property: Applied **independently** to each position in the processed sequence

Applied to Full Sequence

Given sequence $X \in \mathbb{R}^{n \times d_{model}}$ where n is sequence length:

$$\text{FFN}(X) = \begin{bmatrix} \text{FFN}(x_1) \\ \text{FFN}(x_2) \\ \vdots \\ \text{FFN}(x_n) \end{bmatrix}$$

Crucial observation: Same weights W_1, W_2 applied to every position!

```
# Position-wise means:  
for i in range(seq_len):  
    output[i] = FFN(input[i]) # Same FFN for all i
```

Equivalence to 1×1 Convolution

Reshape Perspective

Reshape sequence as "image": $X \in \mathbb{R}^{n \times 1 \times d_{model}}$

- n = "height" (sequence length)
- 1 = "width" (single position)
- d_{model} = "channels"

1×1 Convolution Operation

A 1×1 convolution with kernel $W \in \mathbb{R}^{1 \times 1 \times d_{in} \times d_{out}}$:

$$y_{i,j} = \sum_{c=1}^{d_{in}} W_{c,k} \cdot x_{i,j,c} + b_k$$

For each output channel k , at each spatial position (i, j)

The Mathematical Equivalence

FFN Layer 1: Linear + ReLU

$$Z = \max(0, XW_1 + b_1)$$

As 1×1 conv:

```
Conv1D(kernel_size=1, in_channels=d_model, out_channels=d_ff)
```

Dimension tracking:

```
Input: (batch, seq_len, d_model)  
↓  
↓ reshape for conv  
↓
```

```
(batch, d_model, seq_len) # channels first
↓
↓ 1×1 conv
↓
(batch, d_ff, seq_len)
```

For position i :

$$z_i = \text{ReLU} \left(\sum_{c=1}^{d_{model}} W_1[c, :] \cdot x_i[c] + b_1 \right)$$

This is exactly: **1×1 conv across channel dimension**

FFN Layer 2: Linear

$$\text{Output} = ZW_2 + b_2$$

As 1×1 conv:

```
Conv1D(kernel_size=1, in_channels=d_ff, out_channels=d_model)
```

Complete FFN as Two 1×1 Convolutions

$$\text{FFN}(X) = \text{Conv}_{1\times 1}^{(2)} \left(\text{ReLU} \left(\text{Conv}_{1\times 1}^{(1)}(X) \right) \right)$$

Explicit Form

$$\begin{aligned} \text{FFN}(X) &= \max(0, XW_1 + b_1)W_2 + b_2 \\ &\equiv \text{Conv1D}_{1\times 1}(\text{ReLU}(\text{Conv1D}_{1\times 1}(X))) \end{aligned}$$

PyTorch: Both Implementations are Identical

```
import torch
import torch.nn as nn

# Configuration
batch_size, seq_len, d_model, d_ff = 2, 10, 512, 2048
```

```

# =====
# Implementation 1: Standard FFN (Linear layers)
# =====

class FFN_Linear(nn.Module):
    def __init__(self, d_model, d_ff):
        super().__init__()
        self.linear1 = nn.Linear(d_model, d_ff)
        self.linear2 = nn.Linear(d_ff, d_model)
        self.relu = nn.ReLU()

    def forward(self, x):
        # x: (batch, seq_len, d_model)
        x = self.linear1(x)      # (batch, seq_len, d_ff)
        x = self.relu(x)
        x = self.linear2(x)      # (batch, seq_len, d_model)
        return x

# =====
# Implementation 2: 1x1 Convolution
# =====

class FFN_Conv(nn.Module):
    def __init__(self, d_model, d_ff):
        super().__init__()
        # 1x1 convolution = kernel_size=1
        self.conv1 = nn.Conv1d(d_model, d_ff, kernel_size=1)
        self.conv2 = nn.Conv1d(d_ff, d_model, kernel_size=1)
        self.relu = nn.ReLU()

    def forward(self, x):
        # x: (batch, seq_len, d_model)

        # Conv1d expects (batch, channels, length)
        x = x.transpose(1, 2)    # (batch, d_model, seq_len)

        x = self.conv1(x)        # (batch, d_ff, seq_len)
        x = self.relu(x)
        x = self.conv2(x)        # (batch, d_model, seq_len)

        x = x.transpose(1, 2)    # (batch, seq_len, d_model)
        return x

# =====
# Verify they're equivalent
# =====

ffn_linear = FFN_Linear(d_model, d_ff)
ffn_conv = FFN_Conv(d_model, d_ff)

# Copy weights from linear to conv
ffn_conv.conv1.weight.data = ffn_linear.linear1.weight.data.unsqueeze(2)
ffn_conv.conv1.bias.data = ffn_linear.linear1.bias.data

```

```

ffn_conv.conv2.weight.data = ffn_linear.linear2.weight.data.unsqueeze(2)
ffn_conv.conv2.bias.data = ffn_linear.linear2.bias.data

# Test on same input
x = torch.randn(batch_size, seq_len, d_model)

out_linear = ffn_linear(x)
out_conv = ffn_conv(x)

print(f"Outputs are equal: {torch.allclose(out_linear, out_conv, atol=1e-6)}")
# Output: Outputs are equal: True

print(f"Max difference: {(out_linear - out_conv).abs().max().item()}")
# Output: Max difference: ~1e-7 (numerical precision)

```

Why This Perspective Matters

1. Computational Interpretation

Position-wise = No mixing across positions

```

# FFN does NOT look at neighboring positions
output[i] = f(input[i]) # Only depends on position i

# Unlike attention which mixes:
output[i] = Σj attention[i, j] * input[j] # Depends on all j

```

2. Comparison with CNNs

Operation	Receptive Field	Cross-position Mixing
1×1 Conv (FFN)	Single position	✗ No
3×3 Conv	3 positions	✓ Yes (local)
Attention	All positions	✓ Yes (global)

Insight:

- Attention = Global mixing (routes information)
- FFN = Local processing (transforms at each position)
- Together = Route then compute

3. Parameter Sharing

$$\text{Parameters} = d_{model} \times d_{ff} + d_{ff} \times d_{model}$$

Not $n \times d_{model} \times d_{ff}$ where n is sequence length!

Same weights applied to all positions → **Translation equivariance**

```
# Shift the sequence by 1 position
X_shifted = torch.roll(X, shifts=1, dims=1)

# FFN output is also shifted by 1
assert torch.allclose(
    FFN(X_shifted),
    torch.roll(FFN(X), shifts=1, dims=1)
)
```

4. Why Not Use Larger Kernels?

Could use 3×3, 5×5 convolutions:

```
# This would mix neighboring positions
conv = nn.Conv1d(d_model, d_ff, kernel_size=3)
```

But the Transformer philosophy is to

- Attention handles cross-position mixing (global, content-dependent)
 - FFN handles position-wise transformation (local, position-independent)
 - Clean separation of concerns!
-

Mathematical Equivalence

$$\text{FFN}_{\text{position-wise}} \equiv \text{Conv1D}_{\text{kernel}=1}$$

1. **Conceptual clarity:** FFN doesn't mix positions, only transforms at each position
2. **Computational efficiency:** Can use optimised conv implementations
3. **Framework understanding:** Attention = mixing, FFN = processing
4. **Architecture design:** Could replace with other position-wise operations

In the Bigger Picture

Transformer Layer:

- └─ Attention: Cross-position mixing (global, content-based)
- └─ FFN (1×1 conv): Position-wise processing (local, independent)

Thus Transformers are both powerful and interpretable!

Transformer Decoder Layer

More complex than encoder:

1. Masked Self-Attention:

- Prevents attending to future tokens
- Preserves autoregressive property

2. Cross-Attention:




- Queries from decoder
- Keys and Values from encoder output
- Connects source and target

3. Feed-Forward:

- Same as encoder

Information flow in decoder

Decoder Input



↓

Layer Output

How Encoder's Output is Used in Decoder

The encoder output is used in the **cross-attention** (encoder-decoder attention) sublayer of each decoder layer

- **Encoder output provides Keys (K) and Values (V)**
 - **Decoder provides Queries (Q)**
-

Detailed Architecture Flow

ENCODER (processes source: "The cat sat")

↓
encoder_output = [h₁, h₂, h₃] (one vector per source token)
↓
↓ (This gets fed to EVERY decoder layer)
↓
↓
↓

DECODER (generates target: "Le chat")

Decoder Layer (repeated N times):

1. Masked Self-Attention
(decoder tokens attend to
previous decoder tokens)
Q, K, V all from decoder

↓ (residual + norm)

2. Cross-Attention ← ENCODER HERE!
Q: from decoder current state
K: from encoder_output
V: from encoder_output


Attention(Q_{dec}, K_{enc}, V_{enc})

↓ (residual + norm)

3. Feed-Forward Network

↓ (residual + norm)

Decoder Layer: Information Flow



The Math: How It's Combined

```
# In each decoder layer:  
  
# Step 1: Masked self-attention  
x = decoder_input # Current decoder state  
self_attn_out = MaskedSelfAttention(Q=x, K=x, V=x)  
x = LayerNorm(x + self_attn_out) # Residual connection  
  
# Step 2: Cross-attention (ENCODER OUTPUT USED HERE!)  
cross_attn_out = CrossAttention(  
    Q=x, # Query from current decoder state  
    K=encoder_output, # Keys from encoder ← HERE  
    V=encoder_output # Values from encoder ← HERE
```

```

)
x = LayerNorm(x + cross_attn_out) # Residual connection (ADDED)

# Step 3: Feed-forward
ff_out = FeedForward(x)
x = LayerNorm(x + ff_out) # Residual connection

```

Key

1. Encoder's output is ADDED via residual connection

- just like self-attention output

2. The encoder_output is used as K and V

```

# Cross-attention projects encoder_output
K = encoder_output @ W_K # Keys from source
V = encoder_output @ W_V # Values from source
Q = decoder_state @ W_Q # Query from target

# Then standard attention
attention_weights = softmax(Q @ K.T / sqrt(d_k))
output = attention_weights @ V

```

3. Same encoder_output used in ALL decoder layers

- Each decoder layer has its own cross-attention
- But they all attend to the same encoder_output
 - only the final encoder output is added to each decoder layer
 - there is a key, value pair for each token
 - only some architectures add key-value pair on a encoder-decoder layer-per-layer basis
 - these are not standard Transformers
 - a very high number of parameters, like a dense-net
 - simple and flexible architecture
 - good information flow
 - learning is not incremental
 - layers might focus on different semantic relationships and long-range dependencies
- Different layers learn different attention patterns

4. The encoder_output is a SEQUENCE

```
encoder_output.shape = (batch_size, source_seq_len, d_model)
# One vector for each source token
# Decoder can attend to any/all source positions
```

Complete PyTorch Example

```
class TransformerDecoderLayer(nn.Module):
    def __init__(self, d_model=512, num_heads=8, d_ff=2048, dropout=0.1):
        super().__init__()

        # 1. Masked self-attention
        self.self_attn = MultiHeadAttention(d_model, num_heads)
        self.norm1 = nn.LayerNorm(d_model)
        self.dropout1 = nn.Dropout(dropout)

        # 2. Cross-attention (encoder-decoder attention)
        self.cross_attn = MultiHeadAttention(d_model, num_heads)
        self.norm2 = nn.LayerNorm(d_model)
        self.dropout2 = nn.Dropout(dropout)

        # 3. Feed-forward
        self.ffn = nn.Sequential(
            nn.Linear(d_model, d_ff),
            nn.ReLU(),
            nn.Dropout(dropout),
            nn.Linear(d_ff, d_model)
        )
        self.norm3 = nn.LayerNorm(d_model)
        self.dropout3 = nn.Dropout(dropout)

    def forward(self, x, encoder_output, src_mask=None, tgt_mask=None):
        """
        Args:
            x: Decoder input (batch, target_seq_len, d_model)
            encoder_output: Output from encoder (batch, source_seq_len,
                           d_model)
            src_mask: Mask for source sequence (padding)
            tgt_mask: Causal mask for target sequence
        """

        # 1. Masked self-attention (decoder attends to itself)
        self_attn_out = self.self_attn(
            Q=x, K=x, V=x,
            mask=tgt_mask
        )
        x = self.norm1(x + self.dropout1(self_attn_out)) # Residual +
```

```

norm

    # 2. Cross-attention (decoder attends to encoder)
    cross_attn_out = self.cross_attn(
        Q=x,                                # Query from decoder
        K=encoder_output,                   # Keys from encoder
        V=encoder_output,                   # Values from encoder
        mask=src_mask
    )
    x = self.norm2(x + self.dropout2(cross_attn_out)) # Residual +
norm (ADDED)

    # 3. Feed-forward
    ff_out = self.ffn(x)
    x = self.norm3(x + self.dropout3(ff_out)) # Residual + norm

return x

```

Why This Design?

Query from decoder = "What do I need to generate this target word?"

- Based on what decoder has generated so far
- Changes at each decoding step

Keys & Values from encoder = "Here's what the source sentence contains"

- Fixed encoding of source
- Decoder queries it to find relevant information

Example: Translating "The cat sat" → "Le chat"

When generating "chat":

```

# Decoder state (after generating "Le"):
Q_decoder = "I need to translate the main subject noun"

# Encoder provides:
K_encoder = ["determiner", "ANIMAL", "action", ...] # Keys for matching
V_encoder = [semantic_the, semantic_cat, semantic_sat, ...] # Content

# Cross-attention computes:
attention = softmax(Q_decoder @ K_encoder.T) # High weight on "cat"
output = attention @ V_encoder # Retrieve cat semantics

```

```
# This output is ADDED to decoder state via residual
decoder_state = decoder_state + output # Enriched with source info
```

- Encoder output is used in the **cross-attention sublayer** of each decoder layer
- It provides the **Keys and Values** (decoder provides Query)
- The cross-attention output is **ADDED** to the decoder state via residual connection
- Same encoder output is used by all decoder layers (each learns different attention patterns)

This design **allows the decoder to dynamically "look at" different parts of the source sequence** at each generation step, deciding what source information is relevant for generating each target token.

Masking in Transformer

Padding Mask:

```
# Don't attend to padding tokens
mask = (input_tokens == PAD_TOKEN)
scores = scores.masked_fill(mask, -1e9)
```

Causal Mask (Look-ahead mask):

```
# Prevent attending to future tokens
mask = torch.triu(torch.ones(seq_len, seq_len), diagonal=1).bool()
#   t0  t1  t2  t3
# t0 [0,  1,  1,  1] # Can only see t0
# t1 [0,  0,  1,  1] # Can see t0, t1
# t2 [0,  0,  0,  1] # Can see t0, t1, t2
# t3 [0,  0,  0,  0] # Can see all

scores = scores.masked_fill(mask, -1e9)
attention = softmax(scores) # Masked positions → 0 after softmax
```

7. Positional Encoding

The Position Problem

Attention is permutation-equivariant:

$$\text{Attention}(\pi(X)) = \pi(\text{Attention}(X))$$

for any permutation π

This means:

```
# Without positional info, these are identical:  
["cat", "sat", "mat"] ≡ ["mat", "cat", "sat"]
```

But word order matters!

- "cat sat on mat" \neq "mat sat on cat"
-

Sinusoidal Positional Encoding

Formula:

$$PE_{(pos,2i)} = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

$$PE_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

where:

- pos : position in sequence (0, 1, 2, ...)
- i : dimension index (0, 1, ..., $d_{model}/2$)
- Even dimensions use sine, odd use cosine

Key Properties:

1. **Unique encoding for each position**
 2. **Relative position information:** PE_{pos+k} is a linear function of PE_{pos}
 3. **Bounded values:** $[-1, 1]$
 4. **No learned parameters:** Deterministic function
 5. in original Transformer
-

Adding or concatenating?

1. the new vector may be concatenated

1. this extends the dimensionality of encoding
 2. if one-hot encoding of dimensions, the dimension grows quickly
 3. but is reversible
2. or it may be added
 1. each position is differently coded
 2. irreversible
 3. modified when learning
-

Why Sinusoidal?

Mathematical Insight:

For relative position k :

$$PE(pos + k) = T_k \cdot PE(pos)$$

where T_k is a transformation matrix that depends only on k



$$\sin(\alpha + \beta) = \sin(\alpha) \cos(\beta) + \cos(\alpha) \sin(\beta)$$

This allows model to learn to attend by relative position!

Alternative: Learned Positional Embeddings

```
self.pos_embedding = nn.Embedding(max_seq_len, d_model)
# Trade-off: More flexible but can't extrapolate to longer sequences
```

Positional Encoding in Transformers



2D Positional Encoding

```
def create_2d_position_encoding(h, w, embed_dim):
    """Separate encodings for x and y coordinates"""
    pos_embed_h = sinusoidal_position_encoding(h, embed_dim // 2)
    pos_embed_w = sinusoidal_position_encoding(w, embed_dim // 2)

    # Combine row and column encodings
    pos_embed = torch.cat([
        pos_embed_h.unsqueeze(1).repeat(1, w, 1),
        pos_embed_w.unsqueeze(0).repeat(h, 1, 1)
    ], dim=-1)

    return pos_embed.flatten(0, 1) # (h*w, embed_dim)
```

Relative Position Encoding

Instead of absolute positions, encode relative distances:

$$\text{Attention}_{ij} = \text{softmax} \left(\frac{q_i \cdot k_j + r_{ij}}{\sqrt{d}} \right)$$

where r_{ij} encodes the relative position between patches i and j

Learnable Position Embeddings (ViT Default)

```
self.pos_embed = nn.Parameter(torch.zeros(1, num_patches + 1, embed_dim))
nn.init.trunc_normal_(self.pos_embed, std=0.02)

# Add to patch embeddings
x = x + self.pos_embed
```

Advantages:

- Flexible, can learn any pattern
- Simple to implement
- Empirically effective

Disadvantages:

- Fixed resolution at training
- Doesn't naturally generalize to different image sizes

8. Why Transformers Work: Theoretical Insights

Path Length Between Dependencies

Architecture	Max Path Length	Complexity per Layer	Sequential Ops
RNN	$O(n)$	$O(n \cdot d^2)$	$O(n)$
CNN	$O(\log_k n)$	$O(k \cdot n \cdot d^2)$	$O(1)$
Transformer	$O(1)$	$O(n^2 \cdot d)$	$O(1)$

Key Advantage:

- Any two positions connected in $O(1)$ layers
- Direct gradient paths for all pairs
- Enables learning long-range dependencies

The Transformer as a Composition of Functions

Mathematical View:

A Transformer is:

$$f(x) = f_N \circ f_{N-1} \circ \cdots \circ f_2 \circ f_1(x)$$

where each layer f_i is:

$$f_i(x) = x + \text{FFN}(x + \text{Attention}(x))$$

remember:

- **Attention:** Routes information (soft routing)
 - **FFN:** Processes information (computation)
 - **Residual:** Preserves gradient flow
-

Attention as Soft Dictionary Lookup (Formal)

Theorem: Attention can implement any dictionary lookup

Given keys $K = [k_1, \dots, k_m]$ and values $V = [v_1, \dots, v_m]$:

Hard lookup:

$$\text{lookup}(q) = v_i \quad \text{where } i = \arg \max_j \langle q, k_j \rangle$$

Soft lookup (Attention):

$$\text{Attention}(q) = \sum_j \text{softmax}(\langle q, k_j \rangle / \tau) \cdot v_j$$

As $\tau \rightarrow 0$ (temperature), $\text{Attention}(q) \rightarrow \text{lookup}(q)$

Generalization:

- Multiple queries in parallel: $Q = [q_1, \dots, q_n]$
 - Batch processing all lookups simultaneously
-

Transformer as Universal Approximator

Theorem (Yun et al., 2020):

A Transformer with sufficient depth and width can approximate any sequence-to-sequence function to arbitrary precision.

Key components:

1. **Multi-head attention**: Can implement any sparse connectivity pattern
2. **FFN with ReLU**: Universal function approximation
3. **Depth**: Compositional representations
4. **Residual connections**: Information highways

Intuition:

```
# Can approximate any f: sequence → sequence
y = Transformer(x)

# By composing:
y = FFN_N(Attn_N(...FFN_1(Attn_1(x))))
```

Why Residual Connections Matter

Standard Network:

$$x_{l+1} = f_l(x_l)$$

With Residuals:

$$x_{l+1} = x_l + f_l(x_l)$$

Gradient Flow:

$$\frac{\partial \mathcal{L}}{\partial x_l} = \frac{\partial \mathcal{L}}{\partial x_{l+1}} \left(1 + \frac{\partial f_l}{\partial x_l} \right)$$

The "+1" ensures:

- Gradients can't vanish (always have direct path)
- Network can learn identity if needed: $f_l(x) = 0$
- Enables very deep networks (100+ layers)

Layer Normalisation: Why It's Critical

Formula:

$$\text{LayerNorm}(x) = \gamma \odot \frac{x - \mu}{\sigma + \epsilon} + \beta$$

where

$$\mu = \frac{1}{d} \sum_i x_i$$
$$\sigma^2 = \frac{1}{d} \sum_i (x_i - \mu)^2$$

Why in Transformers?

1. **Stabilises training:** Prevents activation explosion
2. **Enables deeper networks:** Each layer starts with normalised distribution
3. **Reduces dependence on initialisation:** Less sensitive to weight initialisation
4. **Faster convergence:** Smoother loss landscape

Pre-LN vs Post-LN:

```
# Post-LN (original Transformer)
x = LayerNorm(x + Sublayer(x))

# Pre-LN (modern, more stable)
x = x + Sublayer(LayerNorm(x))
```

Modern Transformers predominantly use Pre-LN!

9. Representing Data as Sequences

The Universal Sequence View

Core Idea: Almost any structured data can be represented as a sequence

Text:	"Hello world"	→ [h, e, l, l, o, _, w, o, r, l, d]
Images:	224×224 image	→ 196 patches of 16×16
Graphs:	Social network	→ [node1, node2, ..., nodeN]
Audio:	Waveform	→ [sample1, sample2, ...]
Video:	Frame sequence	→ [frame1, frame2, ...]

Key insight: Sequence processing is a general framework!

Tokenisation: The First Step

What is a Token?

- Atomic unit of input that Transformer processes
- Can be words, subwords, characters, patches, nodes, etc.

Common Tokenization Strategies:

Domain	Tokens	Example
NLP	Words	["The", "cat", "sat"]
NLP	Subwords (BPE)	["The", "cat", "s", "at"]
Vision	Patches	16×16 pixel patches
Audio	Spectrograms	Time-frequency bins
Graphs	Nodes	Graph vertices
Code	Tokens	["def", "function", "(", "x", ")"]

Text as Sequences: NLP

Example: Machine Translation

```
# Tokenization
source = "The cat sat on the mat"
tokens = tokenizer(source) # ["The", "cat", "sat", "on", "the", "mat"]

# Embedding
embeddings = embedding_layer(tokens) # (6, 512)

# Add positional encoding
embeddings = embeddings + positional_encoding

# Process with Transformer
output = transformer(embeddings)

# Decode
translation = decoder(output) # "Le chat s'est assis sur le tapis"
```

Why it works:

- Natural sequential structure

- Order matters (syntax, semantics)
 - Long-range dependencies (anaphora, discourse)
-

Images as Sequences: Vision

Patching Strategy (covered in ViT lecture):

```
# Image: 224x224x3
# Divide into 16x16 patches
# Result: 14x14 = 196 patches

patches = rearrange(image, 'b c (h p1) (w p2) -> b (h w) (p1 p2 c)', 
                     p1=16, p2=16)
# patches: (batch, 196, 768)
```

Sequential View:

Position:	[1]	[2]	[3]	...	[196]
	↓	↓	↓		↓
Patches:	[P ₁]	[P ₂]	[P ₃]	...	[P ₁₉₆]

Each patch is a "word" in the image "sentence"!

Graphs as Sequences: Graph Neural Networks

Challenge: Graphs don't have natural order

Solutions:

1. Arbitrary Ordering:

```
# Order nodes arbitrarily
nodes = [n1, n2, n3, ..., nN]
# Use attention to learn relationships
# Positional encoding less meaningful
```

2. Adjacency-Based Masking:

```
# Only attend to neighbors
attention_mask[i, j] = 1 if edge(i, j) exists
                      0 otherwise
```

3. Graph Transformer:

```
# Structural encoding
structure_encoding = encode_graph_structure(adj_matrix)
node_features = node_features + structure_encoding
output = transformer(node_features)
```

Time Series and Audio

Raw Audio:

```
# Waveform: sequence of samples
audio = [sample_1, sample_2, ..., sample_T]
# Direct application: T can be very large!
```

Spectrogram:

```
# 2D representation: frequency × time
spectrogram = STFT(audio) # (freq_bins, time_frames)
# Treat as sequence of frequency vectors
tokens = [freq_vec_1, freq_vec_2, ..., freq_vec_T]
```

Hierarchical Processing:

```
# Multi-scale: subsample at different rates
coarse = audio[::100] # Every 100th sample
medium = audio[::10] # Every 10th sample
fine = audio # All samples
```

Multimodal: Combining Different Modalities

Vision + Language (e.g., CLIP, DALL-E):

```
# Concatenate sequences from different modalities
image_tokens = vision_encoder(image) # (196, 512)
text_tokens = text_encoder(text) # ( 20, 512)

# Unified sequence
combined = concatenate([image_tokens, text_tokens]) # (216, 512)

# Process with Transformer
```

```

output = transformer(combined)

# Cross-modal attention!
# Image tokens can attend to text tokens and vice versa

```

Key Insight:

- Different modalities → Different tokenizations
 - But same attention mechanism!
 - Shared representation space
-

Why Sequence Representation is Powerful

1. **Unified Framework:** Same architecture for different data types
2. **Flexible Interactions:**
 - Self-attention within modality
 - Cross-attention between modalities
3. **Compositionality:**
 - Complex structures from simple tokens
 - Hierarchical relationships emerge
4. **Scalability:**
 - Parallel processing of all tokens
 - GPU-friendly computation

Trade-off:

- Lose domain-specific inductive biases
 - Need more data to learn structure
 - But gain flexibility and generality!
-

10. Architecture Details

Complete Transformer Encoder

$$Z = \text{LayerNorm}(X + \text{MultiHeadAttention}(X, X, X))$$

$$\text{Output} = \text{LayerNorm}(Z + \text{FFN}(Z))$$

Component Breakdown:

```

class TransformerEncoder(nn.Module):
    def __init__(self, d_model=512, num_heads=8, d_ff=2048, dropout=0.1):
        super().__init__()

        # Multi-head attention
        self.self_attn = MultiHeadAttention(d_model, num_heads)

        # Feed-forward network
        self.ffn = nn.Sequential(
            nn.Linear(d_model, d_ff),
            nn.ReLU(),
            nn.Dropout(dropout),
            nn.Linear(d_ff, d_model)
        )

        # Layer normalization
        self.norm1 = nn.LayerNorm(d_model)
        self.norm2 = nn.LayerNorm(d_model)

        # Dropout
        self.dropout1 = nn.Dropout(dropout)
        self.dropout2 = nn.Dropout(dropout)

    def forward(self, x, mask=None):
        # Self-attention with residual
        attn_output = self.self_attn(x, x, x, mask)
        x = self.norm1(x + self.dropout1(attn_output))

        # Feed-forward with residual
        ffn_output = self.ffn(x)
        x = self.norm2(x + self.dropout2(ffn_output))

    return x

```

[CODE EXAMPLE NEEDED: Full implementation with all components]

Complete Transformer Decoder

```

class TransformerDecoder(nn.Module):
    def __init__(self, d_model=512, num_heads=8, d_ff=2048, dropout=0.1):
        super().__init__()

        # Masked self-attention (causal)
        self.self_attn = MultiHeadAttention(d_model, num_heads)

        # Cross-attention to encoder

```

```

        self.cross_attn = MultiHeadAttention(d_model, num_heads)

        # Feed-forward
        self.ffn = nn.Sequential(
            nn.Linear(d_model, d_ff),
            nn.ReLU(),
            nn.Dropout(dropout),
            nn.Linear(d_ff, d_model)
        )

        # Layer norms
        self.norm1 = nn.LayerNorm(d_model)
        self.norm2 = nn.LayerNorm(d_model)
        self.norm3 = nn.LayerNorm(d_model)

        # Dropout
        self.dropout1 = nn.Dropout(dropout)
        self.dropout2 = nn.Dropout(dropout)
        self.dropout3 = nn.Dropout(dropout)

    def forward(self, x, encoder_output,
               src_mask=None, tgt_mask=None):
        # Masked self-attention
        attn1 = self.self_attn(x, x, x, tgt_mask)
        x = self.norm1(x + self.dropout1(attn1))

        # Cross-attention to encoder
        attn2 = self.cross_attn(x, encoder_output,
                               encoder_output, src_mask)
        x = self.norm2(x + self.dropout2(attn2))

        # Feed-forward
        ffn_out = self.ffn(x)
        x = self.norm3(x + self.dropout3(ffn_out))

    return x

```

Training Details: Teacher Forcing

During Training:

```

# Use ground truth as input (parallel training)
# Input: [SOS, "Le", "chat", "s'est"]
# Target: ["Le", "chat", "s'est", "assis", "EOS"]

for batch in dataloader:

```

```

encoder_out = encoder(source)

# Decoder sees ground truth (shifted right)
decoder_out = decoder(target[:-1], encoder_out)

# Predict next token
loss = criterion(decoder_out, target[1:])

loss.backward()
optimizer.step()

```

During Inference:

```

# Autoregressive generation (sequential)
generated = [SOS]

for _ in range(max_length):
    decoder_out = decoder(generated, encoder_out)
    next_token = argmax(decoder_out[-1])

    if next_token == EOS:
        break

    generated.append(next_token)

```

Practical Hyperparameters

Original Transformer ("Attention is All You Need"):

```

config = {
    'd_model': 512,           # Model dimension
    'num_layers': 6,          # Encoder and decoder layers
    'num_heads': 8,           # Attention heads
    'd_ff': 2048,            # FFN inner dimension
    'dropout': 0.1,           # Dropout rate
    'max_seq_len': 512,       # Maximum sequence length
    'vocab_size': 37000,      # Vocabulary size
}

```

Transformer-Big (Better performance):

```

config_big = {
    'd_model': 1024,
    'num_layers': 6,
    'num_heads': 16,
}

```

```

    'd_ff': 4096,
    'dropout': 0.3,
}

```

Modern Large Models (GPT-3, etc.):

```

config_large = {
    'd_model': 12288,           # 12K dimensions!
    'num_layers': 96,           # 96 layers
    'num_heads': 96,
    'd_ff': 49152,
    'params': '175B',          # 175 billion parameters
}

```

Optimization: Learning Rate Schedule

Warmup + Decay:

```

def get_lr(step, d_model=512, warmup_steps=4000):
    arg1 = step ** (-0.5)
    arg2 = step * (warmup_steps ** (-1.5))
    return (d_model ** (-0.5)) * min(arg1, arg2)

```

$$lr = d_{model}^{-\frac{1}{2}} \cdot \min(\text{step}^{-\frac{1}{2}}, \text{step} \cdot (\text{warmup-steps}^{-\frac{3}{2}}))$$

$$lr = \frac{1}{\sqrt{d_{model}}} \cdot \min \left(\frac{1}{\sqrt{\text{step}}}, \text{step} \cdot \left(\frac{1}{\sqrt{\text{warmup-steps}^3}} \right) \right)$$

Why Warmup?

- Transformers sensitive to initialization
- Large gradients early in training
- Warmup stabilizes training

11. Computational Complexity Analysis

Attention Complexity

Standard Self-Attention:

$$\text{Complexity} = O(n^2 \cdot d)$$

where:

- n : sequence length
- d : model dimension

Breakdown:

1. QK^T : $O(n^2 \cdot d)$ — matrix multiplication
2. Softmax: $O(n^2)$ — row-wise operation
3. Multiply by V : $O(n^2 \cdot d)$

Total: $O(n^2 \cdot d)$ per layer

For full Transformer:

- Encoder: $L \cdot O(n^2 \cdot d)$ where L = number of layers
 - Decoder: Similar but with additional cross-attention
-

Memory Complexity

Storing Attention Matrices:

$$\text{Memory} = O(n^2 \cdot h + n \cdot d)$$

where h is number of heads

Components:

- Attention weights: $h \times n \times n$ (can be large!)
- Activations: $n \times d$ per layer
- Gradients: Same as activations (during training)

Example:

```
# Sequence length n = 1024, d_model = 512, heads = 8
attention_matrix = 8 * 1024 * 1024 * 4 bytes # ~33 MB per layer
# For 12 layers: ~400 MB just for attention!
```

Comparison with Other Architectures

Architecture	Complexity per Layer	Sequential Ops	Max Path Length
RNN	$O(n \cdot d^2)$	$O(n)$	$O(n)$
CNN	$O(k \cdot n \cdot d^2)$	$O(1)$	$O(\log_k(n))$
Attention	$O(n^2 \cdot d)$	$O(1)$	$O(1)$

Trade-offs:

- RNN: Sequential bottleneck, but linear memory
- CNN: Limited receptive field, but efficient
- Transformer: $O(n^2)$ complexity, but parallel + global

Efficient Attention Variants

Problem: $O(n^2)$ doesn't scale to very long sequences

Solutions:

1. Sparse Attention (Routing patterns):

```
# Only attend to local neighbors + few global tokens
attention_pattern = local_window + strided + global_tokens
# Reduces to O(n * k) where k << n
```

2. Linear Attention (Kernelized):

```
# Approximate attention with kernel functions
# O(n) complexity!
attention = φ(Q) @ (φ(K)^T @ V)
# instead of (Q @ K^T) @ V
```

3. Flash Attention (I/O efficient):

```
# Fused CUDA kernels
# Same complexity but much faster in practice
# Reduces memory from O(n^2) to O(n)
```

12. Why Transformers Dominate

Advantages Over Previous Architectures

Aspect	RNN/LSTM	CNN	Transformer
Parallelization	✗ Sequential	✓ Parallel	✓ Parallel
Long-range deps	✗ Gradient vanishing	✗ Limited receptive field	✓ Direct connections
Inductive bias	Strong (sequential)	Strong (locality)	Weak (flexible)
Data efficiency	✓ Good	✓ Good	✗ Needs large data
Interpretability	✗ Hidden state	✗ Feature maps	✓ Attention weights
Scalability	✗ Limited	✗ Saturates	✓ Continues improving

The Scaling Hypothesis

Key Observation: Transformers continue to improve with:

- More parameters
- More data
- More compute

Performance = Power Law(Scale)

$$\log(\text{Loss}) \approx -\alpha * \log(N) + C$$

where N = parameters, data, or compute

This doesn't hold for CNNs/RNNs! They saturate.

Empirical Wins

NLP:

- Machine Translation: BLEU score improvements
- Language Modeling: Perplexity reductions
- Few-shot Learning: Emergent capabilities

Vision:

- ImageNet: State-of-the-art accuracy
- Object Detection: Superior to CNN-based methods
- Video Understanding: Temporal modeling

Multi-modal:

- CLIP: Zero-shot classification
 - DALL-E: Text-to-image generation
 - Flamingo: Few-shot learning across modalities
-

The Lottery: What Makes Transformers Special?

Hypotheses:

1. Inductive Bias Trade-off:

- Weak priors allow learning from data
- Doesn't impose wrong assumptions

2. Expressiveness:

- Universal approximation with fewer constraints
- Can represent more complex functions

3. Optimisation Landscape:

- Residual connections create smooth paths
- Skip connections prevent gradient issues

4. Attention as Routing:

- Soft, learnable connectivity
- Adaptive computation based on content

5. Parallel Training:

- Efficient use of modern hardware
 - Scales better with resources
-

13. Limitations and Challenges

Computational Cost

```
# Example: GPT-3 training
parameters = 175e9
tokens = 300e9
```

```
compute = 3.14e23 FLOPs # ~$4.6M in cloud costs!
training_time = ~1 month on thousands of GPUs
```

Inference Cost:

```
# Single forward pass for GPT-3
sequence_length = 2048
floating_point_ops = 2 * 175e9 * 2048 # ~700 GFLOPs
# Hundreds of ms latency
```

Quadratic Complexity in Sequence Length

Problem:

```
# Memory and compute grow as O(n2)
seq_len = [128, 256, 512, 1024, 2048, 4096]
memory = [x**2 for x in seq_len]
# [16K, 65K, 262K, 1M, 4M, 16M] elements

# Can't process very long sequences!
# Books, long documents, high-res images, videos
```

Partial Solutions:

- Sparse attention patterns
- Linear attention approximations
- Hierarchical processing
- Segmentation strategies

Data Hunger

Observation: Transformers need massive amounts of data

```
# Typical requirements
small_model = {
    'params': '110M',
    'data': '10B tokens', # ~10GB text
    'compute': '1e20 FLOPs'
}

large_model = {
    'params': '175B',
```

```
'data': '300B tokens', # ~300GB text  
'compute': '3e23 FLOPs'  
}
```

Why?

- Weak inductive biases
 - Must learn structure from data
 - Overparameterization requires regularization through data
-

Interpretability Challenges

Attention ≠ Explanation:

```
# Common misconception:  
# "High attention weight means the model uses this information"  
  
# Reality:  
# - Attention is just one component  
# - FFN can override attention  
# - Multiple heads complicate interpretation  
# - Attention can be uniform but still useful
```

Open Questions:

- What do different heads learn?
 - Why do some heads become "no-op"?
 - How does information flow through layers?
 - What concepts do neurons represent?
 - how is knowledge actually stored?
-

14. Modern Variants and Extensions

Encoder-Only Models

BERT (Bidirectional Encoder Representations from Transformers):

```
# Masked language modeling  
input: ["The", "[MASK]", "sat", "on", "the", "mat"]  
output: ["The", "cat", "sat", "on", "the", "mat"]  
  
# Used for:
```

```
# - Classification  
# - Named Entity Recognition  
# - Question Answering
```

Architecture:

- Only encoder stack
 - Bidirectional attention (see full context)
 - Pre-trained on massive corpora
-

Decoder-Only Models

GPT (Generative Pre-trained Transformer):

```
# Autoregressive language modeling  
input: ["The", "cat", "sat"]  
output: "on" # Predict next token  
  
# Used for:  
# - Text generation  
# - Few-shot learning  
# - In-context learning
```

Architecture:

- Only decoder stack (with causal masking)
- Unidirectional attention
- Scales to massive sizes (GPT-3: 175B params)

Key Insight: Decoder-only models can do everything!

- Generation (natural)
 - Classification (with prompting)
 - Translation (with examples)
-

Efficient Transformers

Reformer (Kitaev et al., 2020):

- Locality-sensitive hashing for attention
- Reversible layers (saves memory)
- $O(n \log n)$ complexity

Linformer (Wang et al., 2020):

- Low-rank approximation of attention
- $O(n)$ complexity
- Small accuracy drop

Performer (Choromanski et al., 2021):

- Kernelized attention (FAVOR+)
 - $O(n)$ complexity
 - Maintains performance
-

15. Summary & Key Takeaways

Core Concepts Recap

1. Attention Mechanism:

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

- Soft, differentiable routing of information
- Content-based, not position-based

2. Transformer Architecture:

- Stacked encoder-decoder with attention
- Multi-head attention for diverse patterns
- Position-wise feed-forward networks
- Residual connections + Layer normalization

3. Why It Works:

- $O(1)$ path between any two positions
 - Parallel processing (efficient training)
 - Flexible, learned inductive biases
 - Scales with compute and data
-

When to Use Transformers

Use Transformers When:

- Large dataset available
- Long-range dependencies important

- Parallel training resources available
- Flexibility more important than efficiency
- State-of-the-art performance needed

✗ Consider Alternatives When:

- Small dataset (<10K examples)
 - Real-time inference critical
 - Memory/compute constrained
 - Strong domain priors available
 - Interpretability is paramount
-

The Road Ahead

Current Research Directions:

1. Efficiency:

- Linear attention mechanisms
- Sparse transformers
- Model compression

2. Long Context:

- Extending to 100K+ tokens
- Hierarchical processing
- Memory-augmented transformers

3. Multimodal:

- Unified architectures
- Cross-modal learning
- Few-shot transfer

4. Theory:

- Understanding what is learned
 - Optimization dynamics
 - Generalization bounds
-

16. Practical Tips for Implementation

Debugging Checklist

```
# Common issues and solutions
```

1. **Exploding/Vanishing Gradients**
 - Check: Gradient norms
 - Fix: Gradient clipping, learning rate warmup
2. **Attention Collapse**
 - Check: Attention weight entropy
 - Fix: Dropout, smaller initialization
3. **Training Instability**
 - Check: Loss spikes
 - Fix: Mixed precision, gradient accumulation
4. **Poor Performance**
 - Check: Positional encoding, masking
 - Fix: Verify masks, `try` learned positions
5. **OOM (Out of Memory)**
 - Check: Batch size, sequence length
 - Fix: Gradient checkpointing, smaller batches

Hyperparameter Tuning Guide

Start Here (Defaults):

```
config = {
    'learning_rate': 1e-4,
    'warmup_steps': 4000,
    'batch_size': 32,
    'dropout': 0.1,
    'num_layers': 6,
    'num_heads': 8,
    'd_model': 512,
    'd_ff': 2048,
}
```

Tuning Priority:

1. Learning rate (most impactful)
2. Warmup steps
3. Batch size
4. Dropout
5. Architecture (last resort)

Pre-training vs Fine-tuning

Pre-training (if you have massive compute):

```
# Train from scratch on large corpus
# Requires: 100GB+ data, weeks of training, $$$

model = Transformer(...)
train(model, large_corpus, epochs=100)
```

Fine-tuning (recommended for most):

```
# Start from pre-trained model
# Requires: Small dataset, hours of training, $

model = load_pretrained('bert-base')
model.classifier = nn.Linear(768, num_classes)

# Lower learning rate for fine-tuning
optimizer = AdamW(model.parameters(), lr=2e-5)
train(model, task_data, epochs=3)
```

17. Connections to Future Topics

What's Next in This Course

Week 6: Sequence Models (RNNs, LSTMs, S4)

- How Transformers evolved from RNNs
- Trade-offs between architectures
- When to use what

Week 7: Self-Supervised Learning

- BERT and GPT pre-training objectives
- Contrastive learning with Transformers
- Masked modeling strategies

Week 10: Scaling Laws

- Why Transformers scale so well
- Emergent abilities
- Data/compute trade-offs

Week 12: Mechanistic Interpretability

- What Transformers learn
 - Attention patterns
 - Circuit discovery
-

18. References & Further Reading

Essential Papers

1. [Attention Is All You Need](#) (Vaswani et al., 2017)
 - The original Transformer paper
 - Must-read foundation
 2. [BERT](#) (Devlin et al., 2018)
 - Bidirectional pre-training
 - Encoder-only architecture
 3. [GPT-2](#) and [GPT-3](#) (OpenAI)
 - Decoder-only at scale
 - Few-shot learning emergence
 4. [Vision Transformer \(ViT\)](#) (Dosovitskiy et al., 2020)
 - Covered in Lecture 04
 - Transformers for images
 5. [Formal Algorithms for Transformers](#) (Phuong & Hutter, 2022)
 - Comprehensive mathematical treatment
 - Excellent reference
-

Tutorials and Code

- [The Annotated Transformer](#)
 - Line-by-line implementation
 - Best starting point for coding
- [Hugging Face Transformers](#)
 - Production-ready implementations
 - Pre-trained models
- [PyTorch Transformer Tutorial](#)
 - Official tutorial
 - Clean implementation
- [minGPT](#)

- Minimal GPT implementation
 - By Andrej Karpathy
-

Advanced Topics

- **Efficient Transformers:** [Survey Paper](#)
 - **Scaling Laws:** [Kaplan et al.](#)
 - **Interpretability:** [A Mathematical Framework for Transformer Circuits](#)
 - **Optimization:** [On Layer Normalization](#)
-

Appendix: Mathematical Derivations

A1: Why Scale by $\sqrt{d_k}$?

Problem: Without scaling, dot products grow large in high dimensions

Proof: Assume $q_i, k_j \sim \mathcal{N}(0, 1)$ independently.

$$q \cdot k = \sum_{i=1}^{d_k} q_i k_i$$

$$[q \cdot k] = \sum_{i=1}^{d_k} [q_i][k_i] = 0$$

$$\text{ar}(q \cdot k) = \sum_{i=1}^{d_k} \text{ar}(q_i k_i) = d_k$$

So $q \cdot k \sim \mathcal{N}(0, d_k)$

Issue: Softmax saturates when inputs are large!

$$f(x) : \text{softmax}(x) = [0, \dots, 1, \dots, 0]$$

(near one-hot \rightarrow small gradients)

Solution: Scale by $\sqrt{d_k}$:

$$\frac{q \cdot k}{\sqrt{d_k}} \sim \mathcal{N}(0, 1)$$

Now variance is constant regardless of dimension!

A2: Positional Encoding Properties

Claim: For offset k , PE_{pos+k} is a linear function of PE_{pos}

Proof sketch:

Using trigonometric identities:

$$\sin(\alpha + \beta) = \sin(\alpha) \cos(\beta) + \cos(\alpha) \sin(\beta)$$

$$\cos(\alpha + \beta) = \cos(\alpha) \cos(\beta) - \sin(\alpha) \sin(\beta)$$

We can write:

$$[PE_{pos+k,2i} \ PE_{pos+k,2i+1}] = [\cos(\beta) \ \sin(\beta) \ -\sin(\beta) \ \cos(\beta)] [PE_{pos,2i} \ PE_{pos,2i+1}]$$

$$\text{where } \beta = \frac{k}{10000^{2i/d}}$$

Implication: Relative position k encoded by fixed transformation! Model can learn to attend by relative position.

[DERIVATION NEEDED]: Full mathematical proof]

A3: Multi-Head Attention as Ensemble

Theorem: Multi-head attention with h heads learns h different attention patterns.

Intuition: Each head i has parameters (W_i^Q, W_i^K, W_i^V) that are learned independently.

The attention score for head i :

$$A_i = \text{softmax} \left(\frac{(XW_i^Q)(XW_i^K)^T}{\sqrt{d_k}} \right) (XW_i^V)$$

Final output combines all heads:

$$\text{Output} = \text{Concat}(A_1, \dots, A_h)W^O$$

Each A_i can specialize to different patterns!

Empirical observation:

- Some heads track syntax
 - Some heads track semantics
 - Some heads track positional patterns
 - Redundancy provides robustness
-

[END OF SLIDES]

Next Steps for Students

Before Next Class:

1. **Read:** "Attention Is All You Need" paper
2. **Code:** Implement scaled dot-product attention
3. **Experiment:** Train small Transformer on toy task
4. **Think:** How would you apply Transformers to your research area?

Lab Assignment:

Implement a mini-Transformer for character-level language modeling:

- Dataset: Shakespeare text
- Task: Predict next character
- Model: 4-layer decoder-only Transformer
- Deliverable: Working model + analysis

Project Ideas:

1. Compare Transformer vs RNN on long sequences
 2. Visualize attention patterns in trained model
 3. Implement efficient attention variant
 4. Apply Transformer to new domain (graph, time series, etc.)
-
-