


# 05 Attention and Transformers

## Today's Roadmap

1. **From RNNs to Attention:** The motivation
2. **The Attention Mechanism:** Core concepts and mathematics
3. **Transformer Architecture:** Building blocks
4. **Why Transformers Work:** Theoretical insights
5. **Sequential Data Representation:** Beyond NLP
6. **Architecture Deep Dive:** Implementation details
7. **Computational Considerations:** Efficiency and scaling

## 1. The Problem with Recurrent Architectures

### The Sequential Bottleneck



RNN/LSTM Processing:

$$\begin{array}{ccccccc} h_1 & \rightarrow & h_2 & \rightarrow & h_3 & \rightarrow & h_4 & \rightarrow \dots \rightarrow h_n \\ \uparrow & & \uparrow & & \uparrow & & \uparrow \\ x_1 & & x_2 & & x_3 & & x_4 & & x_n \end{array}$$

Problems:

- Sequential computation (can't parallelize)
- Long-range dependencies vanish
- Information bottleneck through hidden state

## Training RNN

1. training translates to *expansion* of the recursive structure to a multilayer perceptron
2. now we can use gradient descent algorithm

$$L(x_1, \dots, x_T, y_1, \dots, y_T, w_h, w_o) = \frac{1}{T} \sum_{t=1}^T l_t(y_t, o_t)$$

$$\begin{aligned} 3. \quad \frac{\partial L}{\partial w_h} &= \frac{1}{T} \sum_{t=1}^T \frac{\partial l(y_t, o_t)}{\partial w_h} = \frac{1}{T} \sum_{t=1}^T \frac{\partial l(y_t)}{\partial o_t} \frac{\partial g(h_t, w_o)}{\partial h_t} \frac{\partial h_t}{\partial w_h} \\ &\quad \frac{\partial h_t}{\partial w_h} = \frac{\partial f(x_t, h_{t-1}, w_h)}{\partial w_h} + \frac{\partial f(x_t, h_{t-1}, w_h)}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial w_h} \end{aligned}$$

4. last  $\frac{\partial h_t}{\partial w_h}$  (last part) on the weights from the last state  $h_{t-1}$  and, at the same time  $w_h$  (first part after equation)
5. it is possible to make some approximations, but it is still not a stable computation
6. it needs to be computed only on some last  $T$  states only

# The Vanishing Gradient Problem Revisited

Even with LSTM/GRU:


- Path length between distant tokens grows linearly:  $O(n)$
- Gradient flow diminishes over long sequences
- Earlier tokens have exponentially smaller influence

**Mathematical View:**

$$\frac{\partial h_t}{\partial h_0} = \prod_{i=1}^t \frac{\partial h_i}{\partial h_{i-1}}$$


If  $\frac{\partial h_i}{\partial h_{i-1}} < 1$ , gradients vanish as  $t \rightarrow \infty$

- partially solved in recurrent models like **GRU** and **LSTM**
  - introduction of ReLU activation
  - gating



## Gradient Flow: RNN vs Transformer

Recurrent Neural Network (RNN)



- RNN
  - sequential flow
  - diminishing gradient
  - long path with no shortcuts
- Transformer
  - flow by layers with residual connections

- stable gradients with slight degradation

## The Fixed Context Problem

**Encoder-Decoder RNN:**

```
# Entire input compressed into single vector c
encoder_output = encode(x1, x2, ..., xn) # → c (fixed size!)
decoder_output = decode(c, y1, y2, ..., ym)

# Problem: c must capture EVERYTHING about the input
# This is an information bottleneck!
```

"rnn-encoder-decoder.svg" could not be found.

**Example: Translation**

English: "The agreement on the European Economic Area was signed in August 1992"  
[14 words → compressed to fixed vector c → decode]

For each output word, decoder sees:

- Same context vector c
- No direct access to specific input words
- Can't "look back" at relevant parts

## 2. Attention

"Instead of encoding the entire input into a fixed context vector, let the decoder **attend** to different parts of the input at each decoding step."

### The Problem: Fixed Context Bottleneck

- before attention
  - encoder-decoder compressed the **entire input sequence into a single fixed-size context vector\*\***
  - e.g. "*The agreement on the European Economic Area was signed in August 1992*"
  - all 14 words had to be squeezed into one vector
  - severe information bottleneck, especially for long sequences
  - loosing important relations between words

### The Bahdanau Innovation

[Bahdanau et al. \(2014\)](#) (over 40 thousand citations) first attention mechanism for neural machine translation.

- instead of using a fixed context vector,
- decoder **dynamically attends to different parts of the encoder output** at each decoding step:

1. **Encoder produces a sequence of hidden states:**  $h_1, h_2, \dots, h_n$  (one per input word)
2. **At each decoder step t,** compute **alignment scores** between decoder state  $s_{t-1}$  and each encoder state

$$h_j : e_{tj} = \text{score}(s_{t-1}, h_j) = v^T \tanh(W_1 s_t + W_2 h_j)$$

3. **Convert to attention weights** via softmax:


$$\alpha_{tj} = \frac{\exp(e_{tj})}{\sum_{k=1}^n \exp(e_{tk})}$$

#### 4. Compute context vector as weighted sum:

$$c_t = \sum_{j=1}^n \alpha_{tj} h_j$$

5. Use  $c_t$  for decoding : Generate output word based  $s_t$  on both decoder state and context

6. It is possible to build two series of hidden states:  $h_1, h_2, \dots, h_n$  (from first to last word) and  $h_n, h_{n-1}, \dots, h_1$  (from last to first) and combine them



#### Example in translation:

Translating: "The cat sat" → "Le chat s'est assis"

When generating "chat" (French for cat):

- Decoder computes attention over ["The", "cat", "sat"]
- High attention weight on "cat":  $\alpha = [0.1, 0.8, 0.1]$
- Context vector emphasizes "cat" representation
- Decoder uses this to generate correct translation

## Model training

1. specify the encoder and decoder models (LSTM, GRU)
2. add attention module
3. build a translation module (sequence-to-sequence model too)
4. train all

## Why It Mattered

### Impact on the field:

- **Solved the bottleneck:** No need to compress entire sequence into single vector
- **Enabled long sequences:** Could handle 50+ word sentences effectively
- **Interpretable:** Attention weights showed which source words influenced each target word

- **State-of-the-art results:** Dramatically improved translation quality

### Difference from Transformers:

- **Still used RNNs:** Attention was an *add-on* to RNN encoder-decoder
- **Sequential processing:** Had to process words one at a time (no parallelization)
- **Auxiliary mechanism:** Attention helped RNNs, but RNNs were still the core architecture

### The path to Transformers:

```
Bahdanau Attention (2014): "Let's add attention TO RNNs"
    ↓
    Huge improvement!
    ↓
Vaswani et al. (2017): "What if we use ONLY attention?"
    ↓
    Transformer is born
```

## Mathematical Formulation

The Bahdanau attention mechanism can be written as:

$$\text{context}_t = \sum_{j=1}^n \text{softmax(score}(s_t, h_j)) \cdot h_j$$

where the score function learns to measure alignment between decoder state  $s_t$  and encoder state  $h_j$

**This is conceptually similar to the Transformer's attention:**

$$\text{Attention}(Q, K, V) = \text{softmax}(QK^T)V$$

but

- Bahdanau: Single query (decoder state) attending to sequence
- Transformer: All positions attend to all positions simultaneously
- Bahdanau: Score computed via learned MLP
- Transformer: Score computed via dot product (simpler, faster)

### Bahdanau attention

- proved that explicit attention mechanisms could dramatically improve sequence-to-sequence models.
- showed that models could "look back" at relevant input positions rather than relying on a compressed context
- inspired the Transformer, which took the idea further by making attention the *only* mechanism, removing RNNs entirely.

### [Transformer explained](#)

## Attention Intuition

Translation: "The cat sat on the mat" → "Le chat s'est assis sur le tapis"

When generating "chat":

Attention weights: [ 0.8, 0.15, 0.02, 0.01, 0.01, 0.01 ]  
 ↓      ↓      ↓      ↓      ↓      ↓  
 chat    s'est    assis    sur    le    tapis

Model focuses on "cat" (0.8 weight) while translating!

## From Implicit to Explicit Attention

- **RNN Encoder-Decoder (2014):**
  - Implicit attention through recurrent connections
    - Each  $h_t$  depends on all previous  
 $\$\$ h_t = f(h_{t-1}, x_t) \$\$$
- **Bahdanau Attention (2015):**
  - Explicit attention weights

$$\alpha_t = \text{softmax}(\text{score}(h_t, \text{encoder-states}))$$

$$\text{context}_t = \sum \alpha_t[i] * \text{encoder-states}[i]$$


- **Self-Attention / Transformer (2017):**
  - Attention as the PRIMARY mechanism (not auxiliary)
  - No RNN at all

$$\text{output} = \text{Attention}(Q, K, V)$$

## 5. Attention Mechanism

- model needs to cope with input passages (e.g. in NLP) of different lengths
- know connections between elements that depend on attention
- transformer gets both by using **dot-product self-attention**

### Self-Attention: The Core Operation



(from Understanding deep learning)

- standard network computes a linear transformation with a non-linear function
- **self-attention  $sa[]$**  takes
  - $N$  inputs
  - returns  $N$  output vectors of the same size (e.g., in NLP a word or some sub-word)
- a **value** is computed from each token

$$v_m = \beta_v + \Omega_v x_m$$


using biases  $\beta$  and weights  $\Omega$

- may be done in parallel
- a scalar  $a[x_i, x_j]$  is the **attention** that token  $x_j$  pays to token  $x_i$ 
  - attentions  $a[\cdot, x_j]$  sum up to 1
- **self attention**

$$sa_j(x_1, \dots, x_N) = \sum_{i=1}^N a[x_i, x_j] \cdot v_i$$

- a weighted sum of all values  $v_i$

- weights  $a(\cdot, x_n)$  are non-negative and sum-up to 1
- each self-attention  $sa_i[x_1, \dots]$  can be thought as routing of the original  $N$  tokens with different proportions for the current task
- all can be computed in parallel per token
- in the figure above
  - $N$  inputs are taken
  - in the left-most a routing to  $sa_1(x)$  is computed with weights 0.1, 0.3, 0.6
  - then two different routings



(from Prince, Understanding deep learning, MIT, 2023)

- query vectors are computed as  $q_n = \beta_q + \Omega_q x_n$
- key vectors are computed as  $k_n = \beta_k + \Omega_k x_n$
- dot products are passed to a softmax giving attention values

Given input  $\mathbf{X} \in \mathbb{R}^{N \times D}$ :

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Where:

- $Q = XW_Q$  (Queries): "What am I looking for?"
- $K = XW_K$  (Keys): "What information do I have?"
- $V = XW_V$  (Values): "What information to aggregate?"
- where dimensions are
  - $d$  model dimension, size of the embedding,
  - $d_k$  the key and query dimension:  $d_k = d/n_{heads}$  (standard relationship)
  - $d_q = d_k$  (typically  $d_q = d_k$ )
  - $d_v$  dimension of values in attention:  $d_v = d$ 
    - typically embedding dimension needs to be a multiple of the number of heads


- if  $d_q \neq d_k$ , an additional attention may be used where additional trained matrices map both to the same dimension

$$W_Q \in \mathbb{R}^{d \times d_q}, W_K \in \mathbb{R}^{d \times d_k}, W_V \in \mathbb{R}^{d \times d_v}$$

- dot products in computation may get to large values
- and the softmax regions where largest value dominates, and the gradients get very small
- model gets hard to train
- the scaling by the square root of the number of rows of keys (and queries) prevents it

## Key, query, value


- the **dot product** returns a measure of similarity between its arguments
- weights  $a[x_., x_n]$  depend on relative similarities between the n-th query and all the keys
- queries and keys should have the same dimensions
  - there is a possibility of mapping to a common value




---

## Multiple attention heads

- multiple heads may be computed in parallel



- typically for model dimension  $d$  and  $num_{heads}$  heads, the values, queries and keys will **all be of the same size** allowing for efficient computation
- multiple heads are concatenated

### 3. Understanding Query, Key, and Value: The Heart of Attention

#### Beyond the Matrices: Building Intuition

[Transformer explainer](#)

#### The Fundamental Question

Why do we need THREE separate matrices (Q, K, V)?

**Bad answer:** "Because the math works out"

- in my (/igor) opinion the [Attention Is All You Need, Vaswani et al., 2017](#) paper caused Transformers a great deal of harm
- people now understand Transformer just as clever multiplication of matrices
- almost nobody asks themselves how the model "thinks", stores knowledge

**Good answer:** "Because attention is answering three distinct questions:"

- Query (Q):** "What information am I looking for?"
- Key (K):** "What information do I offer?"
- Value (V):** "What is the actual information I provide?"

## Library Search

```
You walk into a library looking for:  
QUERY: "Books about neural networks written after 2017"
```

### The librarian's process:

#### 1. Check catalog cards (KEYS):

- Book 1: "Deep Learning, 2016, Goodfellow"
- Book 2: "Neural Networks, 2020, Smith"
- Book 3: "Transformers, 2017, Vaswani"

#### 2. Match your query against keys:

- Book 1: X (too old)
- Book 2: ✓ ✓ (perfect match!)
- Book 3: ✓ (borderline - exactly 2017)

#### 3. Retrieve actual books (VALUES):

- The catalog card (KEY) tells you WHERE to look
- But what you get is the actual BOOK (VALUE)
- Key = "metadata for matching"
- Value = "the content you actually want"

### In attention:

```
Your query: "Books about neural networks after 2017"  
↓  
Similarity scores: [0.1, 0.9, 0.5] # How well each book matches  
↓  
Weighted retrieval: 0.1 * Book1 + 0.9 * Book2 + 0.5 * Book3  
↓  
You get: Mostly Book 2, some Book 3, almost no Book 1
```

## Database Query (SQL-like)

```
SELECT value_column  
FROM table  
WHERE key_column MATCHES query  
ORDER BY similarity(key_column, query) DESC  
  
-- In attention, this becomes:  
-- Weighted retrieval based on soft matching!
```

### Example: Student Database

```
# You want information about: "CS students with GPA > 3.5"  
query = [field: "CS", criteria: "GPA > 3.5"]  
  
# Database entries:  
students = [  
    {key: "Alice, CS, 3.9",      value: "Research: NLP, Skills: PyTorch"},  
    {key: "Bob, Math, 3.8",     value: "Research: Topology, Skills: Proofs"},  
    {key: "Carol, CS, 3.2",     value: "Research: Graphics, Skills: OpenGL"},  
    {key: "Dave, CS, 3.7",      value: "Research: ML, Skills: TensorFlow"}  
]  
  
# Attention scores (how well keys match query):  
scores = [0.95, 0.1, 0.3, 0.9] # Alice and Dave match best  
  
# Weighted retrieval (soft SQL):  
result = 0.95 * "Research: NLP..." + 0.1 * "Research: Topology..."  
       + 0.3 * "Research: Graphics..." + 0.9 * "Research: ML..."
```

```
# You get: Mixture of Alice and Dave's info (CS students with high GPA)
```

### Key insight:

- **Keys** are for matching/indexing
- **Values** are the actual content you retrieve
- **Query** is what you're searching for

---

## Search Engine

```
User types: "best restaurants near me" ← QUERY  
↓
```

Google's index:

```
Document 1:  
KEY: [location:nearby, topic:restaurants,  
      rating:4.5, recency:2024]  
VALUE: "Mario's Pizza: 123 Main St,  
      4.5★, Italian cuisine..."
```

```
Document 2:  
KEY: [location:nearby, topic:hotels,  
      rating:4.0, recency:2023]  
VALUE: "Grand Hotel: 456 Oak Ave,  
      4.0★, Luxury accommodation..."
```

### Matching process:

1. **Encode query:** "restaurants nearby" → query vector
2. **Compare to keys:** Which documents are relevant?
  - Doc 1: topic=restaurants ✓, location=nearby ✓ → HIGH score
  - Doc 2: topic=hotels X, location=nearby ✓ → LOW score
3. **Retrieve values:** Return actual document content (weighted by scores)

### Why separate Key and Value?

- **Key:** Optimised for fast similarity search
  - Compressed representation: [location, topic, rating, recency]
  - Like a hash/index for quick lookup
- **Value:** The full, rich information
  - Complete document content
  - Much larger, more detailed
  - You don't search through this directly (too expensive!)

---

## The Mathematical Picture: Information Retrieval

### From Retrieval to Continuous Attention

#### Hard retrieval (traditional):

```
def hard_lookup(query, keys, values):  
    # Find BEST matching key  
    best_idx = argmax(similarity(query, keys))  
    # Return corresponding value  
    return values[best_idx]  
  
# Problem: Discrete, non-differentiable
```

## Soft retrieval (attention):

```
def soft_lookup(query, keys, values):
    # Compute similarity to ALL keys
    scores = similarity(query, keys) # [s1, s2, ..., sn]

    # Convert to probabilities (softmax)
    weights = softmax(scores) # [w1, w2, ..., wn]

    # Weighted average of ALL values
    return Σ wi * values[i]

# Differentiable! Can learn what to attend to
```

- most Transformer architectures use some form of top-k architecture
  - in NN top-k means applying the weights modification **only** to the highest **k** outputs
  - thus only the meaningful classes are modified
  - it can be done using differentiable operations, not just cutting off

## Visualisation of scores → weights:

```
Raw scores:      [2.1,      0.3,      -0.5,      1.8]
                  ↓ softmax
Attention weights: [0.58,      0.10,      0.04,      0.28] # Sum to 1.0
                    ↓
Weighted sum:    0.58*V1 + 0.10*V2 + 0.04*V3 + 0.28*V4
```

## Translating "The cat sat on the mat" → French

### Setup: Decoder generating "chat" (French for "cat")

#### Semantic dimensions:

1. **Is-Animal** (0-1)
2. **Is-Action** (0-1)
3. **Is-Object** (0-1)
4. **Definiteness** (0-1)

#### Step 1: The Query (from decoder)

Decoder's current state: "I'm trying to generate the French word for 'cat'"

```
Q = [0.9, 0.1, 0.2, 0.3]
      |   |   |
      |   |   |   └ definiteness (the cat, not just any cat)
      |   |   |   └ is-object (0.2 - can be object but not here)
      |   |   |   └ is-action (0.1 - not an action)
      |   |   |   └ is-animal (0.9 - STRONGLY looking for animal!)

# Query says: "I need something that is primarily an ANIMAL,
#               not an action, possibly definite"
```

#### Step 2: The Keys (from encoder - what each word "advertises")

#### Each English word broadcasts what it contains:

```
# "The" (first occurrence)
K_the1 = [0.0, 0.0, 0.0, 1.0]
          |   |   |
          |   |   |   └ definiteness=1.0 (it's "the"!)
          |   |   |   └ is-animal=0.0 (not an animal)

# "cat"
```

```

K_cat = [1.0, 0.0, 0.3, 0.0]
      |           |           └ definiteness=0.0 (not a determiner)
      |           |           └ is-object=0.3 (can be object)
      |           |           └ is-action=0.0 (not an action)
      |           |           └ is-animal=1.0 (YES! I'm an animal!)

# "sat"
K_sat = [0.0, 1.0, 0.0, 0.0]
      |           └ is-action=1.0 (I'm a verb/action!)
      |           └ is-animal=0.0 (not an animal)

# "on"
K_on = [0.0, 0.0, 0.0, 0.0]
      └ (preposition - low on all semantic features)

# "the" (second occurrence)
K_the1 = [0.0, 0.0, 0.0, 1.0]
      └ definiteness=1.0

# "mat"
K_mat = [0.0, 0.0, 1.0, 0.0]
      |           └ is-object=1.0 (I'm an object!)
      |           └ is-action=0.0
      |           └ is-animal=0.0 (not an animal)

```

### Step 3: Compute Attention Scores ( $\mathbf{Q} \cdot \mathbf{K}$ )

```
Q = [is-animal, is-action, is-object, definiteness]
```

Element-wise multiplication and sum:

```

# Q · K_the1
score_the1 = 0.9×0.0 + 0.1×0.0 + 0.2×0.0 + 0.3×1.0 = 0.3
              └ no animal match └ some definiteness

# Q · K_cat
score_cat = 0.9×1.0 + 0.1×0.0 + 0.2×0.3 + 0.3×0.0 = 0.96
              └ STRONG animal match └ minor object match

# Q · K_sat
score_sat = 0.9×0.0 + 0.1×1.0 + 0.2×0.0 + 0.3×0.0 = 0.1
              └ action doesn't match our query

# Q · K_on
score_on = 0.9×0.0 + 0.1×0.0 + 0.2×0.0 + 0.3×0.0 = 0.0
              └ preposition, no match

# Q · K_the2
score_the2 = 0.9×0.0 + 0.1×0.0 + 0.2×0.0 + 0.3×1.0 = 0.3
              └ same as first "the"

# Q · K_mat
score_mat = 0.9×0.0 + 0.1×0.0 + 0.2×1.0 + 0.3×0.0 = 0.2
              └ some object match

```

Raw scores: [0.30, 0.96, 0.10, 0.00, 0.30, 0.20]

	the <sub>1</sub>	cat	sat	on	the <sub>2</sub>	mat
	[0.30]	0.96	0.10	0.00	0.30	0.20

### Step 4: Scale by $\sqrt{d_k}$

```

d_k = 4 # dimension of keys; needed for scalability
scaling_factor = √4 = 2.0

scaled_scores = [0.30, 0.96, 0.10, 0.00, 0.30, 0.20] / 2.0
= [0.15, 0.48, 0.05, 0.00, 0.15, 0.10]

```

## Step 5: Apply Softmax → Attention Weights

```

# exp(scaled_scores)
exp_scores = [exp(0.15), exp(0.48), exp(0.05), exp(0.00), exp(0.15), exp(0.10)]
= [1.16, 1.62, 1.05, 1.00, 1.16, 1.11]

# Normalize (sum = 7.10)
attention_weights = [1.16, 1.62, 1.05, 1.00, 1.16, 1.11] / 7.10
= [0.16, 0.23, 0.15, 0.14, 0.16, 0.16]
α ≈ [0.16, 0.23, 0.15, 0.14, 0.16, 0.16]
the1 CAT sat on the2 mat

```

### Visualization:

Word:	the <sub>1</sub>	cat	sat	on	the <sub>2</sub>	mat
Weight:	16%	23%	15%	14%	16%	16%
	—	—	—	—	—	—
	↑					
	Highest attention to "cat"!					

## Step 6: The Values (actual semantic content to retrieve)

Values contain rich, contextualized information:

```

# V_the1: Determiner introducing "cat"
V_the1 = [0.1, 0.0, 0.0, 0.8]
    |                               ↳ definite article marker
    |                               ↳ minimal semantic content

# V_cat: Rich animal semantics
V_cat = [0.9, 0.0, 0.1, 0.7]
    |   |   |   ↳ definite (preceded by "the")
    |   |   |   ↳ minor object role
    |   |   |   ↳ not action
    |   |   |   ↳ ANIMAL (feline, domestic, pet)

# V_sat: Action/state information
V_sat = [0.0, 0.9, 0.0, 0.3]
    |   |   ↳ PAST ACTION (sitting)
    |   |   ↳ not animal

# V_on: Prepositional relationship
V_on = [0.0, 0.0, 0.5, 0.0]
    |   |   ↳ spatial relationship
    |   |   ↳ not animal

# V_the2: Determiner for "mat"
V_the2 = [0.0, 0.0, 0.0, 0.9]
    |   |   ↳ definite article

# V_mat: Object semantics
V_mat = [0.0, 0.0, 0.9, 0.6]
    |   |   |   ↳ definite object
    |   |   |   ↳ OBJECT (floor covering)
    |   |   |   ↳ not animal

```

### **Step 7: Weighted Sum of Values (Final Output)**

```

output = α₁·V_the₁ + α₂·V_cat + α₃·V_sat + α₄·V_on + α₅·V_the₂ + α₆·V_mat

= 0.16·[0.1, 0.0, 0.0, 0.8] + 0.23·[0.9, 0.0, 0.1, 0.7]
+ 0.15·[0.0, 0.9, 0.0, 0.3] + 0.14·[0.0, 0.0, 0.5, 0.0]
+ 0.16·[0.0, 0.0, 0.0, 0.9] + 0.16·[0.0, 0.0, 0.9, 0.6]

# Dimension 1 (is-animal):
= 0.16×0.1 + 0.23×0.9 + 0.15×0.0 + 0.14×0.0 + 0.16×0.0 + 0.16×0.0
= 0.016 + 0.207 + 0 + 0 + 0 + 0
= 0.223

# Dimension 2 (is-action):
= 0.16×0.0 + 0.23×0.0 + 0.15×0.9 + 0.14×0.0 + 0.16×0.0 + 0.16×0.0
= 0 + 0 + 0.135 + 0 + 0 + 0
= 0.135

# Dimension 3 (is-object):
= 0.16×0.0 + 0.23×0.1 + 0.15×0.0 + 0.14×0.5 + 0.16×0.0 + 0.16×0.9
= 0 + 0.023 + 0 + 0.070 + 0 + 0.144
= 0.237

# Dimension 4 (definiteness):
= 0.16×0.8 + 0.23×0.7 + 0.15×0.3 + 0.14×0.0 + 0.16×0.9 + 0.16×0.6
= 0.128 + 0.161 + 0.045 + 0 + 0.144 + 0.096
= 0.574

output = [0.223,      0.135,      0.237,      0.574]
          |           |           |           |
          |           |           |           |   ↳ Strong definiteness signal
          |           |           |           |
          |           |           |           |   ↳ Some object context
          |           |           |           |
          |           |           |           |   ↳ Some action context
          |           |           |           |
          |           |           |           |   ↳ Animal signal (from "cat")

```

## Interpretation of the Output

The enriched representation [0.22, 0.14, 0.24, 0.57] means:

```
Original query was looking for: ANIMAL

Got back:
• 22% animal features      ← from "cat" (23%)
• 14% action features      ← from "sat" (15%)
• 24% object features      ← from "mat" (16%)
• 57% definiteness         ← from "the"s

Decoder now knows:
"Generate an animal word (cat → chat),
it's definite (le/la),
with some action/object context"
```

### **Why this helps translation:**

- Primary signal: ANIMAL (0.22) → Generate "chat" (cat)
  - Definiteness (0.57) → Use "le chat" not just "chat"
  - Context from action/object → Past tense, spatial relationship

## Visual Summary: The Complete Flow

INPUT: "The cat sat on the mat"  
↓  
ENCODER REPRESENTATIONS:

Word	Key (for matching)	Value (content)
the <sub>1</sub>	[0.0, 0.0, 0.0, 1.0]	[0.1, 0.0, 0.0, 0.8]
cat	[1.0, 0.0, 0.3, 0.0]	[0.9, 0.0, 0.1, 0.7]
sat	[0.0, 1.0, 0.0, 0.0]	[0.0, 0.9, 0.0, 0.3]
on	[0.0, 0.0, 0.0, 0.0]	[0.0, 0.0, 0.5, 0.0]
the <sub>2</sub>	[0.0, 0.0, 0.0, 1.0]	[0.0, 0.0, 0.0, 0.9]
mat	[0.0, 0.0, 1.0, 0.0]	[0.0, 0.0, 0.9, 0.6]

↓

DECODER QUERY: [0.9, 0.1, 0.2, 0.3]  
 "Looking for an animal"

↓

ATTENTION SCORES (Q·K):  
 [0.30, 0.96, 0.10, 0.00, 0.30, 0.20]  
 ↑↑↑ Highest match!

↓

ATTENTION WEIGHTS (softmax):  
 [0.16, 0.23, 0.15, 0.14, 0.16, 0.16]  
 ↑↑↑ Focus on "cat"

↓

OUTPUT (weighted sum of Values): \_\_\_\_\_  
 [0.22, 0.14, 0.24, 0.57]  
 └ Mostly animal features, definite, some context

↓

DECODER GENERATES: "chat"

## 1. Keys are for matching

- K\_cat = [1.0, 0.0, 0.3, 0.0] says "I'm an animal"
- Query Q = [0.9, 0.1, 0.2, 0.3] matches strongly (score = 0.96)
- Other keys don't match as well

## 2. Values are for content

- Even though we matched on "cat" via Key
- We retrieve rich Value: [0.9, 0.0, 0.1, 0.7]
- Contains more than just "is-animal" flag
- Has definiteness, context, semantic richness

## 3. Soft retrieval averages context

- Not just "cat" (23% weight)
- Also gets definiteness from "the" (16% + 16%)
- Some action context from "sat" (15%)
- Some object context from "mat" (16%)

## 4. Different queries would give different results

- If query was [0.1, 0.9, 0.1, 0.1] (looking for action)
  - Would attend to "sat" instead
  - Would retrieve action semantics
  - Would generate French verb

## Why Not Just Use One Matrix?

### Thought Experiment: What if Q = K = V?

```
# Self-attention with Q = K = V = X
attention = softmax(X @ X.T) @ X
```

```
# This means:  
# - Matching criterion = Content itself  
# - Retrieved content = Same as matching criterion
```

### Problem 1: Conflation of "what to match" and "what to retrieve"

Example from translation:

Word "bank" in: "I went to the bank by the river"

If Key = Value:

- Key says: "I'm the word 'bank' (financial OR river-side)"
- Value says: Same information

But we want:

- Key to say: "I'm a noun that relates to 'river'"  
→ Helps matching in this context
- Value to provide: Full semantic content for "river bank"  
→ Different from just the word form

### Problem 2: Reduced expressiveness

```
# With Q = K = V (1 matrix):
```

Parameters:  $d^2$

```
# With Q, K, V separate (3 matrices):
```

Parameters:  $3d^2$

# More parameters = more expressive power

# Can learn richer representations

### Problem 3: Asymmetry of roles

Query: "What do I need?" (generated by current state)

Key: "What can I offer?" (advertised by candidates)

Value: "Here's what I actually give" (delivered content)

These are fundamentally different questions!

Forcing them to be the same limits what the model can learn.

---

## The Key-Value Separation: A Deeper Dive

### Why is Key ≠ Value Important?

Intuition: Index vs Content

Think of a book:

- **Key** = Chapter titles, keywords, page numbers (the INDEX)
  - Optimized for: Quick scanning, pattern matching
  - Characteristics: Compressed, abstract, matchable
- **Value** = Actual chapter content (the TEXT)
  - Optimized for: Rich information, semantic content
  - Characteristics: Detailed, specific, informative

```
# Book: "Deep Learning"  
key_chapter3 = "Chapter 3: Linear Algebra Review"  
            # Concise, tells you what's inside  
  
value_chapter3 = ""  
    Linear algebra is the study of vectors, matrices...  
    [20 pages of detailed content]
```

```
# Rich, detailed information
```

## In neural networks:


```
# Encoding a sentence: "The black cat"

# Keys: Compressed, abstract representations for matching
K_the = embed("determiner, definite")
K_black = embed("adjective, color, descriptor")
K_cat = embed("noun, animal, subject")

# Values: Rich semantic embeddings with full context
V_the = embed("article with context from 'black cat'")
V_black = embed("color=black, modifies=cat, visual_property")
V_cat = embed("animal=feline, color=black, subject=yes, ...")
    # Much richer! Includes compositional information
```

## Key vs Value: Index vs Content

Why we separate what we match on (Key) from what we retrieve (Value)



## 4. The Attention Mechanism: Mathematics

### Core Attention Equation

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Where:

- $Q$  (Query): "What am I looking for?" —  $\mathbb{R}^{n \times d_k}$  (no of queries × key / query dimension)
- $K$  (Key): "What do I contain?" —  $\mathbb{R}^{m \times d_k}$  (no of keys / values × key / query dimension)
- $V$  (Value): "What information do I have?" —  $\mathbb{R}^{m \times d_v}$  (no of keys / values × value dimension)

### Dimensions:

- $n$ : Number of queries (target sequence length)
- $m$ : Number of keys/values (source sequence length)
- $d_k$ : Key/query dimension (for scalability - see below)
- $d_v$ : Value dimension

---

## Attention as Soft Dictionary Lookup

Analogy:

```
# Traditional dictionary
dictionary = {
    "cat": "a small domesticated carnivorous mammal",
    "dog": "a domesticated carnivorous mammal",
    ...
}
result = dictionary["cat"] # Hard lookup: exact match

# Attention as soft dictionary
queries = ["ct", "catt", "dog"] # Fuzzy queries
keys = ["cat", "dog", "bird"]
values = [embedding_cat, embedding_dog, embedding_bird]

# Each query attends to ALL keys, weighted by similarity
result = Σ similarity(query, key_i) * value_i
```

---

## Step-by-Step Computation

Given input sequence  $X = [x_1, x_2, \dots, x_n] \in \mathbb{R}^{n \times d}$ :

### Step 1: Linear Projections

$$Q = XW^Q, \quad K = XW^K, \quad V = XW^V$$

where  $W^Q, W^K \in \mathbb{R}^{d \times d_k}$  and  $W^V \in \mathbb{R}^{d \times d_v}$

### Step 2: Compute Attention Scores

$$\text{scores} = QK^T \in \mathbb{R}^{n \times n}$$

Element  $(i, j)$  measures how much query  $i$  should attend to key  $j$

### Step 3: Scale (Important!)

$$\text{scaled scores} = \frac{QK^T}{\sqrt{d_k}}$$

Why divide by  $\sqrt{d_k}$ ?

$$q \cdot k = \sum_i^d q_i k_i$$

- $q_i$  and  $k_i$  have 0 mean and 1 variance
- $q_i k_i$  has variance 1 too
- the sum's variance is  $d$  and standard deviation  $\sqrt{d}$
- if scaled  $q \cdot k / \sqrt{d}$ , the whole variance is 1 too regardless of dimension
- division keeps gradients stable (prevents softmax saturation)

---

### Step 4: Apply Softmax

$$\alpha = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right)$$

Each row sums to 1:  $\sum_j \alpha_{ij} = 1$

## Step 5: Weighted Sum of Values

$$\text{Output} = \alpha V$$

---

## Attention Score Interpretation

$$\text{score}_{ij} = q_i \cdot k_j = |q_i| |k_j| \cos(\theta_{ij})$$

- High score → query  $i$  and key  $j$  are "aligned"
- Low score → query  $i$  doesn't need key  $j$

After softmax:

$$\alpha_{ij} = \frac{\exp(\text{score}_{ij})}{\sum_k \exp(\text{score}_{ik})}$$

Attention weights form a probability distribution over source positions

---

## Self-Attention vs Cross-Attention

Type	Q from	K,V from	Use Case
Self-Attention	Same sequence	Same sequence	Encoding contextual relationships
Cross-Attention	Target sequence	Source sequence	Encoder-Decoder connection

### Self-Attention:

```
# All positions attend to all other positions in same sequence
Q = K = V = X # Same source
# Learns internal structure and dependencies
```

### Cross-Attention:

```
# Decoder attends to encoder
Q = decoder_states # What decoder is looking for
K = V = encoder_states # Information from encoder
# Connects source and target sequences
```

---

## 5. Multi-Head Attention

### The Motivation

Single attention head:

- Learns one notion of "relatedness"
- May miss different types of relationships

### Example:

```
Sentence: "The bank by the river is steep"
```

```
Head 1 might learn: syntactic relationships (subject-verb)
Head 2 might learn: semantic relationships (bank-river)
Head 3 might learn: positional relationships (nearby words)
```

## Multi-Head Attention Formula

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$

where each head:

$$\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$$

where

- $h$ : number of heads (typically 8 or 16)
  - $W_i^Q, W_i^K \in \mathbb{R}^{d_{\text{model}} \times d_k}$  where  $d_k = d_{\text{model}}/h$
  - $W_i^V \in \mathbb{R}^{d_{\text{model}} \times d_v}$  where  $d_v = d_{\text{model}}/h$
  - $W^O \in \mathbb{R}^{hd_v \times d_{\text{model}}}$
- 

## Why Multiple Heads?

### Mathematical Intuition:

Instead of one  $d_{\text{model}}$ -dimensional attention:

$$\text{Attention}(Q, K, V) \in \mathbb{R}^{n \times d_{\text{model}}}$$

Split into  $h$  heads of dimension  $d_k = d_{\text{model}}/h$ :

$$\text{head}_i \in \mathbb{R}^{n \times d_k}$$

1. **Subspace specialization:** Each head can attend to different aspects
2. **Parameter efficiency:**  $h$  small projections vs 1 large projection
3. **Ensemble effect:** Multiple attention patterns averaged

Different attention patterns across heads [Transformer explained](#)

---

## Computational View

```
# Pseudocode for Multi-Head Attention

def multi_head_attention(Q, K, V, num_heads=8):
    d_k = d_model // num_heads

    # Split into heads
    Q_heads = split_heads(Q, num_heads)  # (batch, heads, seq, d_k)
    K_heads = split_heads(K, num_heads)
    V_heads = split_heads(V, num_heads)

    # Parallel attention for each head
    attention_outputs = []
    for i in range(num_heads):
        head_i = scaled_dot_product_attention(
            Q_heads[:, :, i], K_heads[:, :, i], V_heads[:, :, i])
        attention_outputs.append(head_i)

    # Concatenate and project
    concat = concatenate(attention_outputs)  # (batch, seq, d_model)
    output = linear(concat, W_0)

    return output
```

[Multihead attention in PyTorch](#)

## Additional attention

If queries and keys do not have the same dimension, we may use an **additional** attention


$$attn(q, k) = w_v \tanh(W_q q + W_k k)$$


## 6. The Transformer Architecture

Three types of transformer architecture

- **encoder** transforms the embeddings into some representation that would support some processing task
- **decoder** predicts the next token to continue the input
- **encoder-decoder** for conversion of one sequence into another e.g., translation

### High-Level View






(after Vaswani et al.)

Full Transformer encoder (left) and decoder (right) architecture (one layer both)

- masked multi-head attention masks future tokens preventing glimpses "to the future"



(after Anthropic blog)

## Transformer Encoder Layer

$$\text{EncoderLayer}(X) = \text{FFN}(\text{LayerNorm}(X + \text{MultiHead}(X, X, X)))$$

**Detailed Steps:**

### 1. Self-Attention with Residual:

$$Z = X + \text{MultiHead}(X, X, X)$$

### 2. Layer Normalization:

$$Z' = \text{LayerNorm}(Z)$$

### 3. Feed-Forward with Residual:

$$H = Z' + \text{FFN}(Z')$$

### 4. Layer Normalisation:

$$\text{Output} = \text{LayerNorm}(H)$$

## Feed-Forward Network (FFN)

### Architecture:

$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$

### Properties:

- Applied **position-wise** (same FFN for each position independently)
- Two linear transformations with ReLU activation
- Typical dimensions:  $d_{model} = 512$ ,  $d_{ff} = 2048$

### Why FFN after Attention?

1. **Non-linearity**: Attention is linear operations + softmax
2. **Mixing information**: Attention aggregates, FFN processes
3. **Capacity**: Adds parameters for complex transformations

## FFN as 1x1 Convolution

### The Position-Wise Feed-Forward Network

#### Standard FFN Definition

$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$

where:

- $x \in \mathbb{R}^{d_{model}}$  is a single position's representation
- $W_1 \in \mathbb{R}^{d_{model} \times d_{ff}}$ , typically  $d_{ff} = 4 \cdot d_{model}$
- $W_2 \in \mathbb{R}^{d_{ff} \times d_{model}}$  squeezes back to  $d_{model}$

**Key property:** Applied **independently** to each position in the processed sequence

### Applied to Full Sequence

Given sequence  $X \in \mathbb{R}^{n \times d_{model}}$  where  $n$  is sequence length:

$$\text{FFN}(X) = \begin{bmatrix} \text{FFN}(x_1) \\ \text{FFN}(x_2) \\ \vdots \\ \text{FFN}(x_n) \end{bmatrix}$$

**Crucial** : Same weights  $W_1, W_2$  applied to every position!

## Equivalence to $1 \times 1$ Convolution

### Reshape Perspective

Reshape sequence as "image":  $X \in \mathbb{R}^{n \times 1 \times d_{model}}$

- $n$  = "height" (sequence length)
- $1$  = "width" (single position)
- $d_{model}$  = "channels"

### $1 \times 1$ Convolution Operation

A  $1 \times 1$  convolution with kernel  $W \in \mathbb{R}^{1 \times 1 \times d_{in} \times d_{out}}$ :

$$y_{i,j} = \sum_{c=1}^{d_{in}} W_{c,k} \cdot x_{i,j,c} + b_k$$

For each output channel  $k$ , at each spatial position  $(i, j)$

## The Mathematical Equivalence

### FFN Layer 1: Linear + ReLU

$$Z = \max(0, XW_1 + b_1)$$

As  $1 \times 1$  conv:

```
Conv1D(kernel_size=1, in_channels=d_model, out_channels=d_ff)
```

Dimension tracking:

```
Input: (batch, seq_len, d_model)
      ↓
      ↓ reshape for conv
      ↓
      (batch, d_model, seq_len) # channels first
      ↓
      ↓ 1×1 conv
      ↓
      (batch, d_ff, seq_len)
```

For position  $i$ :

$$z_i = \text{ReLU} \left( \sum_{c=1}^{d_{model}} W_1[c, :] \cdot x_i[c] + b_1 \right)$$

This is exactly:  **$1 \times 1$  conv across channel dimension**

### FFN Layer 2: Linear

$$\text{Output} = ZW_2 + b_2$$

As  $1 \times 1$  conv:

```
Conv1D(kernel_size=1, in_channels=d_ff, out_channels=d_model)
```

## Complete FFN as Two $1 \times 1$ Convolutions

$$\text{FFN}(X) = \text{Conv}_{1 \times 1}^{(2)} \left( \text{ReLU} \left( \text{Conv}_{1 \times 1}^{(1)}(X) \right) \right)$$

## Explicit Form

$\text{FFN}(X) = \max(0, XW_1 + b_1)W_2 + b_2$ $\equiv \text{Conv1D}_{1 \times 1}(\text{ReLU}(\text{Conv1D}_{1 \times 1}(X)))$
--

---

## PyTorch: Both Implementations are Identical

```

import torch
import torch.nn as nn

# Configuration
batch_size, seq_len, d_model, d_ff = 2, 10, 512, 2048

# =====
# Implementation 1: Standard FFN (Linear layers)
# =====
class FFN_Linear(nn.Module):
    def __init__(self, d_model, d_ff):
        super().__init__()
        self.linear1 = nn.Linear(d_model, d_ff)
        self.linear2 = nn.Linear(d_ff, d_model)
        self.relu = nn.ReLU()

    def forward(self, x):
        # x: (batch, seq_len, d_model)
        x = self.linear1(x)      # (batch, seq_len, d_ff)
        x = self.relu(x)
        x = self.linear2(x)      # (batch, seq_len, d_model)
        return x

# =====
# Implementation 2: 1x1 Convolution
# =====
class FFN_Conv(nn.Module):
    def __init__(self, d_model, d_ff):
        super().__init__()
        # 1x1 convolution = kernel_size=1
        self.conv1 = nn.Conv1d(d_model, d_ff, kernel_size=1)
        self.conv2 = nn.Conv1d(d_ff, d_model, kernel_size=1)
        self.relu = nn.ReLU()

    def forward(self, x):
        # Conv1d expects (batch, channels, length)
        x = x.transpose(1, 2)    # (batch, d_model, seq_len)

        x = self.conv1(x)        # (batch, d_ff, seq_len)
        x = self.relu(x)
        x = self.conv2(x)        # (batch, d_model, seq_len)

        x = x.transpose(1, 2)    # (batch, seq_len, d_model)
        return x

# =====
# Verify they're equivalent
# =====
ffn_linear = FFN_Linear(d_model, d_ff)
ffn_conv = FFN_Conv(d_model, d_ff)

# Copy weights from linear to conv
ffn_conv.conv1.weight.data = ffn_linear.linear1.weight.data.unsqueeze(2)
ffn_conv.conv1.bias.data = ffn_linear.linear1.bias.data
ffn_conv.conv2.weight.data = ffn_linear.linear2.weight.data.unsqueeze(2)
ffn_conv.conv2.bias.data = ffn_linear.linear2.bias.data

```

```

# Test on same input
x = torch.randn(batch_size, seq_len, d_model)

out_linear = ffn_linear(x)
out_conv = ffn_conv(x)

print(f"Outputs are equal: {torch.allclose(out_linear, out_conv, atol=1e-6)}")
# Output: Outputs are equal: True

print(f"Max difference: {(out_linear - out_conv).abs().max().item()}")
# Output: Max difference: ~1e-7 (numerical precision)

```

---

## Why This Perspective Matters

### 1. Computational Interpretation

#### Position-wise = No mixing across positions

```

# FFN does NOT look at neighboring positions
output[i] = f(input[i]) # Only depends on position i

# Unlike attention which mixes:
output[i] = Σj attention[i, j] * input[j] # Depends on all j

```

### 2. Comparison with CNNs

Operation	Receptive Field	Cross-position Mixing
1×1 Conv (FFN)	Single position	✗ No
3×3 Conv	3 positions	✓ Yes (local)
Attention	All positions	✓ Yes (global)

#### Insight:

- Attention = Global mixing (routes information)
  - FFN = Local processing (transforms at each position)
  - Together = Route then compute
- 

### 3. Parameter Sharing

$$\text{Parameters} = d_{model} \times d_{ff} + d_{ff} \times d_{model}$$

Not  $n \times d_{model} \times d_{ff}$  where  $n$  is sequence length!

Same weights applied to all positions → **Translation equivariance**

```

# Shift the sequence by 1 position
X_shifted = torch.roll(X, shifts=1, dims=1)

# FFN output is also shifted by 1
assert torch.allclose(
    FFN(X_shifted),
    torch.roll(FFN(X), shifts=1, dims=1)
)

```

---

### 4. Why Not Use Larger Kernels?

Could use 3×3, 5×5 convolutions:

```
# This would mix neighboring positions
conv = nn.Conv1d(d_model, d_ff, kernel_size=3)
```

### But the Transformer philosophy is to

- Attention handles cross-position mixing (global, content-dependent)
- FFN handles position-wise transformation (local, position-independent)
- Clean separation of concerns!

### Mathematical Equivalence

$$\text{FFN}_{\text{position-wise}} \equiv \text{Conv1D}_{\text{kernel}=1}$$

1. **Conceptual clarity:** FFN doesn't mix positions, only transforms at each position
2. **Computational efficiency:** Can use optimised conv implementations
3. **Framework understanding:** Attention = mixing, FFN = processing
4. **Architecture design:** Could replace with other position-wise operations

### In the Bigger Picture

Transformer Layer:

- |— Attention: Cross-position mixing (global, content-based)
- |— FFN (1×1 conv): Position-wise processing (local, independent)

Thus Transformers are both powerful and interpretable!

## Transformer Decoder Layer


More complex than encoder:

1. **Masked Self-Attention:**
  - Prevents attending to future tokens
  - Preserves autoregressive property
2. **Cross-Attention:**
  - Queries from decoder
  - Keys and Values from encoder output
  - Connects source and target
3. **Feed-Forward:**
  - Same as encoder


### Information flow in decoder

Decoder Input


↓



↓



↓



```
| + Add & Norm |
```

↓  
Layer Output

## How Encoder's Output is Used in Decoder

The encoder output is used in the **cross-attention** (encoder-decoder attention) sublayer of each decoder layer

- **Encoder output provides Keys (K) and Values (V)**
- **Decoder provides Queries (Q)**

## Detailed Architecture Flow

```
ENCODER (processes source: "The cat sat")
↓
encoder_output = [h1, h2, h3] (one vector per source token)
↓
↓ (This gets fed to EVERY decoder layer)
↓
↓
DECODER (generates target: "Le chat")
```

Decoder Layer (repeated N times):

```
1. Masked Self-Attention
   (decoder tokens attend to
    previous decoder tokens)
   Q, K, V all from decoder
```

↓ (residual + norm)

```
2. Cross-Attention ← ENCODER HERE!
   Q: from decoder current state
   K: from encoder_output
   V: from encoder_output
```


Attention(Q<sub>dec</sub>, K<sub>enc</sub>, V<sub>enc</sub>)

↓ (residual + norm)

```
3. Feed-Forward Network
```

↓ (residual + norm)

## Decoder Layer: Information Flow



## The Math: How It's Combined

```
# In each decoder layer:

# Step 1: Masked self-attention
x = decoder_input # Current decoder state
self_attn_out = MaskedSelfAttention(Q=x, K=x, V=x)
x = LayerNorm(x + self_attn_out) # Residual connection

# Step 2: Cross-attention (ENCODER OUTPUT USED HERE!)
cross_attn_out = CrossAttention(
    Q=x, # Query from current decoder state
    K=encoder_output, # Keys from encoder ← HERE
    V=encoder_output # Values from encoder ← HERE
)
x = LayerNorm(x + cross_attn_out) # Residual connection (ADDED)

# Step 3: Feed-forward
ff_out = FeedForward(x)
x = LayerNorm(x + ff_out) # Residual connection
```

## Key

### 1. Encoder's output is ADDED via residual connection

- just like self-attention output

### 2. The encoder\_output is used as K and V

```
# Cross-attention projects encoder_output
K = encoder_output @ W_K # Keys from source
V = encoder_output @ W_V # Values from source
```

```

Q = decoder_state @ W_Q    # Query from target

# Then standard attention
attention_weights = softmax(Q @ K.T / sqrt(d_k))
output = attention_weights @ V

```

### 3. Same encoder\_output used in ALL decoder layers

- Each decoder layer has its own cross-attention
- But they all attend to the same encoder\_output
  - only the final encoder output is added to each decoder layer
  - there is a key, value pair for each token
  - only some architectures add key-value pair on a encoder-decoder layer-per-layer basis
    - these are not standard Transformers
    - a very high number of parameters, like a dense-net
  - simple and flexible architecture
  - good information flow
  - learning is not incremental
  - layers might focus on different semantic relationships and long-range dependencies
- Different layers learn different attention patterns

### 4. The encoder\_output is a SEQUENCE

```

encoder_output.shape = (batch_size, source_seq_len, d_model)
# One vector for each source token
# Decoder can attend to any/all source positions

```

## Complete PyTorch Example

```

class TransformerDecoderLayer(nn.Module):
    def __init__(self, d_model=512, num_heads=8, d_ff=2048, dropout=0.1):
        super().__init__()

        # 1. Masked self-attention
        self.self_attn = MultiHeadAttention(d_model, num_heads)
        self.norm1 = nn.LayerNorm(d_model)
        self.dropout1 = nn.Dropout(dropout)

        # 2. Cross-attention (encoder-decoder attention)
        self.cross_attn = MultiHeadAttention(d_model, num_heads)
        self.norm2 = nn.LayerNorm(d_model)
        self.dropout2 = nn.Dropout(dropout)

        # 3. Feed-forward
        self.ffn = nn.Sequential(
            nn.Linear(d_model, d_ff),
            nn.ReLU(),
            nn.Dropout(dropout),
            nn.Linear(d_ff, d_model)
        )
        self.norm3 = nn.LayerNorm(d_model)
        self.dropout3 = nn.Dropout(dropout)

    def forward(self, x, encoder_output, src_mask=None, tgt_mask=None):
        ....
        Args:
            x: Decoder input (batch, target_seq_len, d_model)
            encoder_output: Output from encoder (batch, source_seq_len, d_model)
            src_mask: Mask for source sequence (padding)
            tgt_mask: Causal mask for target sequence
        ....
        # 1. Masked self-attention (decoder attends to itself)
        self_attn_out = self.self_attn(

```

```

        Q=x, K=x, V=x,
        mask=tgt_mask
    )
x = self.norm1(x + self.dropout1(self_attn_out)) # Residual + norm

# 2. Cross-attention (decoder attends to encoder)
cross_attn_out = self.cross_attn(
    Q=x, # Query from decoder
    K=encoder_output, # Keys from encoder
    V=encoder_output, # Values from encoder
    mask=src_mask
)
x = self.norm2(x + self.dropout2(cross_attn_out)) # Residual + norm (ADDED)

# 3. Feed-forward
ff_out = self.ffn(x)
x = self.norm3(x + self.dropout3(ff_out)) # Residual + norm

return x

```

## Why This Design?

**Query from decoder = "What do I need to generate this target word?"**

- Based on what decoder has generated so far
- Changes at each decoding step

**Keys & Values from encoder = "Here's what the source sentence contains"**

- Fixed encoding of source
- Decoder queries it to find relevant information

**Example: Translating "The cat sat" → "Le chat"**

When generating "chat":

```

# Decoder state (after generating "Le"):
Q_decoder = "I need to translate the main subject noun"

# Encoder provides:
K_encoder = ["determiner", "ANIMAL", "action", ...] # Keys for matching
V_encoder = [semantic_the, semantic_cat, semantic_sat, ...] # Content

# Cross-attention computes:
attention = softmax(Q_decoder @ K_encoder.T) # High weight on "cat"
output = attention @ V_encoder # Retrieve cat semantics

# This output is ADDED to decoder state via residual
decoder_state = decoder_state + output # Enriched with source info

```

- Encoder output is used in the **cross-attention sublayer** of each decoder layer
- It provides the **Keys and Values** (decoder provides Query)
- The cross-attention output is **ADDED** to the decoder state via residual connection
- Same encoder output is used by all decoder layers (each learns different attention patterns)

This design **allows the decoder to dynamically "look at" different parts of the source sequence** at each generation step, deciding what source information is relevant for generating each target token.

## Masking in Transformer

**Padding Mask:**

```
# Don't attend to padding tokens
mask = (input_tokens == PAD_TOKEN)
scores = scores.masked_fill(mask, -1e9)
```

#### Causal Mask (Look-ahead mask):

```
# Prevent attending to future tokens
mask = torch.triu(torch.ones(seq_len, seq_len), diagonal=1).bool()
# t0 t1 t2 t3
# t0 [0, 1, 1, 1] # Can only see t0
# t1 [0, 0, 1, 1] # Can see t0, t1
# t2 [0, 0, 0, 1] # Can see t0, t1, t2
# t3 [0, 0, 0, 0] # Can see all

scores = scores.masked_fill(mask, -1e9)
attention = softmax(scores) # Masked positions → 0 after softmax
```

## 7. Positional Encoding

### The Position Problem

Attention is permutation-equivariant:

$$\text{Attention}(\pi(X)) = \pi(\text{Attention}(X))$$

for any permutation  $\pi$

This means:

```
# Without positional info, these are identical:
["cat", "sat", "mat"] ≈ ["mat", "cat", "sat"]
```

But word order matters!

- "cat sat on mat" ≠ "mat sat on cat"

### Sinusoidal Positional Encoding

Formula:

$$PE_{(pos,2i)} = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

$$PE_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

where:

- $pos$ : position in sequence (0, 1, 2, ...)
- $i$ : dimension index (0, 1, ...,  $d_{model}/2$ )
- Even dimensions use sine, odd use cosine

Key Properties:

1. Unique encoding for each position
2. Relative position information:  $PE_{pos+k}$  is a linear function of  $PE_{pos}$
3. Bounded values:  $[-1, 1]$
4. No learned parameters: Deterministic function
5. in original Transformer

## Adding or concatenating?

1. the new vector may be concatenated
  1. this extends the dimensionality of encoding
  2. if one-hot encoding of dimensions, the dimension grows quickly
  3. but is reversible
2. or it may be added
  1. each position is differently coded
  2. irreversible
  3. modified when learning

## Why Sinusoidal?

### Mathematical Insight:

For relative position  $k$ :

$$PE(pos + k) = T_k \cdot PE(pos)$$

where  $T_k$  is a transformation matrix that depends only on  $k$


$$\sin(\alpha + \beta) = \sin(\alpha) \cos(\beta) + \cos(\alpha) \sin(\beta)$$

This allows model to learn to attend by relative position!


### Alternative: Learned Positional Embeddings

```
self.pos_embedding = nn.Embedding(max_seq_len, d_model)
# Trade-off: More flexible but can't extrapolate to longer sequences
```

## Positional Encoding in Transformers



## Summary: Why Positional Encoding Works



## 2D Positional Encoding

```

def create_2d_position_encoding(h, w, embed_dim):
    """Separate encodings for x and y coordinates"""
    pos_embed_h = sinusoidal_position_encoding(h, embed_dim // 2)
    pos_embed_w = sinusoidal_position_encoding(w, embed_dim // 2)

    # Combine row and column encodings
    pos_embed = torch.cat([
        pos_embed_h.unsqueeze(1).repeat(1, w, 1),
        pos_embed_w.unsqueeze(0).repeat(h, 1, 1)
    ], dim=-1)

    return pos_embed.flatten(0, 1) # (h*w, embed_dim)

```

## Relative Position Encoding

Instead of absolute positions, encode relative distances:

$$\text{Attention}_{ij} = \text{softmax} \left( \frac{q_i \cdot k_j + r_{ij}}{\sqrt{d}} \right)$$

where  $r_{ij}$  encodes the relative position between patches  $i$  and  $j$

## Learnable Position Embeddings (ViT Default)

```

self.pos_embed = nn.Parameter(torch.zeros(1, num_patches + 1, embed_dim))
nn.init.trunc_normal_(self.pos_embed, std=0.02)

# Add to patch embeddings
x = x + self.pos_embed

```

### Advantages:

- Flexible, can learn any pattern
- Simple to implement
- Empirically effective

### Disadvantages:

- Fixed resolution at training
- Doesn't naturally generalize to different image sizes

## 8. Why Transformers Work: Theoretical Insights

### Path Length Between Dependencies

Architecture	Max Path Length	Complexity per Layer	Sequential Ops
RNN	$O(n)$	$O(n \cdot d^2)$	$O(n)$
CNN	$O(\log_k n)$	$O(k \cdot n \cdot d^2)$	$O(1)$
Transformer	$O(1)$	$O(n^2 \cdot d)$	$O(1)$

### Key Advantage:

- Any two positions connected in  $O(1)$  layers
- Direct gradient paths for all pairs
- Enables learning long-range dependencies

# The Transformer as a Composition of Functions

## Mathematical View:

A Transformer is:

$$f(x) = f_N \circ f_{N-1} \circ \cdots \circ f_2 \circ f_1(x)$$

where each layer  $f_i$  is:

$$f_i(x) = x + \text{FFN}(x + \text{Attention}(x))$$

remember:

- **Attention:** Routes information (soft routing)
- **FFN:** Processes information (computation)
- **Residual:** Preserves gradient flow

---

## Attention as Soft Dictionary Lookup (Formal)

**Theorem:** Attention can implement any dictionary lookup

Given keys  $K = [k_1, \dots, k_m]$  and values  $V = [v_1, \dots, v_m]$ :

**Hard lookup:**

$$\text{lookup}(q) = v_i \quad \text{where } i = \arg \max_j \langle q, k_j \rangle$$

**Soft lookup (Attention):**

$$\text{Attention}(q) = \sum_j \text{softmax}(\langle q, k_j \rangle / \tau) \cdot v_j$$

As  $\tau \rightarrow 0$  (temperature),  $\text{Attention}(q) \rightarrow \text{lookup}(q)$

**Generalization:**

- Multiple queries in parallel:  $Q = [q_1, \dots, q_n]$
- Batch processing all lookups simultaneously

---

## Transformer as Universal Approximator

**Theorem** (Yun et al., 2020):

A Transformer with sufficient depth and width can approximate any sequence-to-sequence function to arbitrary precision.

**Key components:**

1. **Multi-head attention:** Can implement any sparse connectivity pattern
2. **FFN with ReLU:** Universal function approximation
3. **Depth:** Compositional representations
4. **Residual connections:** Information highways

**Intuition:**

```
# Can approximate any f: sequence → sequence
y = Transformer(x)

# By composing:
y = FFN_N(Attn_N(...FFN_1(Attn_1(x))))
```

## Why Residual Connections Matter

Standard Network:

$$x_{l+1} = f_l(x_l)$$

With Residuals:

$$x_{l+1} = x_l + f_l(x_l)$$

Gradient Flow:

$$\frac{\partial \mathcal{L}}{\partial x_l} = \frac{\partial \mathcal{L}}{\partial x_{l+1}} \left( 1 + \frac{\partial f_l}{\partial x_l} \right)$$

The "+1" ensures:

- Gradients can't vanish (always have direct path)
- Network can learn identity if needed:  $f_l(x) = 0$
- Enables very deep networks (100+ layers)

## Layer Normalisation: Why It's Critical

Formula:

$$\text{LayerNorm}(x) = \gamma \odot \frac{x - \mu}{\sigma + \epsilon} + \beta$$

where

$$\begin{aligned}\mu &= \frac{1}{d} \sum_i x_i \\ \sigma^2 &= \frac{1}{d} \sum_i (x_i - \mu)^2\end{aligned}$$

Why in Transformers?

1. **Stabilises training:** Prevents activation explosion
2. **Enables deeper networks:** Each layer starts with normalised distribution
3. **Reduces dependence on initialisation:** Less sensitive to weight initialisation
4. **Faster convergence:** Smoother loss landscape

Pre-LN vs Post-LN:

```
# Post-LN (original Transformer)
x = LayerNorm(x + Sublayer(x))

# Pre-LN (modern, more stable)
x = x + Sublayer(LayerNorm(x))
```

Modern Transformers predominantly use Pre-LN!

## 9. Representing Data as Sequences

### The Universal Sequence View

Core Idea: Almost any structured data can be represented as a sequence

Text:	"Hello world"	→ [h, e, l, l, o, _, w, o, r, l, d]
Images:	224×224 image	→ 196 patches of 16×16
Graphs:	Social network	→ [node1, node2, ..., nodeN]

Audio:	Waveform	→ [sample1, sample2, ...]
Video:	Frame sequence	→ [frame1, frame2, ...]

**Key insight:** Sequence processing is a general framework!

---

## Tokenisation: The First Step

### What is a Token?

- Atomic unit of input that Transformer processes
- Can be words, subwords, characters, patches, nodes, etc.

### Common Tokenization Strategies:

Domain	Tokens	Example
NLP	Words	["The", "cat", "sat"]
NLP	Subwords (BPE)	["The", "cat", "s", "at"]
Vision	Patches	16×16 pixel patches
Audio	Spectrograms	Time-frequency bins
Graphs	Nodes	Graph vertices
Code	Tokens	["def", "function", "(", "x", ")"]

---

## Text as Sequences: NLP

### Example: Machine Translation

```
# Tokenization
source = "The cat sat on the mat"
tokens = tokenizer(source) # ["The", "cat", "sat", "on", "the", "mat"]

# Embedding
embeddings = embedding_layer(tokens) # (6, 512)

# Add positional encoding
embeddings = embeddings + positional_encoding

# Process with Transformer
output = transformer(embeddings)

# Decode
translation = decoder(output) # "Le chat s'est assis sur le tapis"
```

### Why it works:

- Natural sequential structure
  - Order matters (syntax, semantics)
  - Long-range dependencies (anaphora, discourse)
- 

## Images as Sequences: Vision

### Patching Strategy (covered in ViT lecture):

```
# Image: 224×224×3
# Divide into 16×16 patches
# Result: 14×14 = 196 patches

patches = rearrange(image, 'b c (h p1) (w p2) -> b (h w) (p1 p2 c)',
```

```
p1=16, p2=16)
# patches: (batch, 196, 768)
```

### Sequential View:

```
Position: [1] [2] [3] ... [196]
          ↓   ↓   ↓   ↓
Patches: [P1] [P2] [P3] ... [P196]
```

Each patch is a "word" in the image "sentence"!

---

## Graphs as Sequences: Graph Neural Networks

**Challenge:** Graphs don't have natural order

### Solutions:

#### 1. Arbitrary Ordering:

```
# Order nodes arbitrarily
nodes = [n1, n2, n3, ..., nN]
# Use attention to learn relationships
# Positional encoding less meaningful
```

#### 2. Adjacency-Based Masking:

```
# Only attend to neighbors
attention_mask[i, j] = 1 if edge(i, j) exists
                      0 otherwise
```

#### 3. Graph Transformer:

```
# Structural encoding
structure_encoding = encode_graph_structure(adj_matrix)
node_features = node_features + structure_encoding
output = transformer(node_features)
```

---

## Time Series and Audio

### Raw Audio:

```
# Waveform: sequence of samples
audio = [sample_1, sample_2, ..., sample_T]
# Direct application: T can be very large!
```

### Spectrogram:

```
# 2D representation: frequency × time
spectrogram = STFT(audio) # (freq_bins, time_frames)
# Treat as sequence of frequency vectors
tokens = [freq_vec_1, freq_vec_2, ..., freq_vec_T]
```

### Hierarchical Processing:

```
# Multi-scale: subsample at different rates
coarse = audio[::100] # Every 100th sample
medium = audio[::10] # Every 10th sample
fine = audio # All samples
```

---

## Multimodal: Combining Different Modalities

**Vision + Language** (e.g., CLIP, DALL-E):

```
# Concatenate sequences from different modalities
image_tokens = vision_encoder(image)      # (196, 512)
text_tokens = text_encoder(text)          # ( 20, 512)

# Unified sequence
combined = concatenate([image_tokens, text_tokens]) # (216, 512)

# Process with Transformer
output = transformer(combined)

# Cross-modal attention!
# Image tokens can attend to text tokens and vice versa
```

**Key Insight:**

- Different modalities → Different tokenizations
- But same attention mechanism!
- Shared representation space

---

## Why Sequence Representation is Powerful

1. **Unified Framework:** Same architecture for different data types
2. **Flexible Interactions:**
  - Self-attention within modality
  - Cross-attention between modalities
3. **Compositionality:**
  - Complex structures from simple tokens
  - Hierarchical relationships emerge
4. **Scalability:**
  - Parallel processing of all tokens
  - GPU-friendly computation

**Trade-off:**

- Lose domain-specific inductive biases
- Need more data to learn structure
- But gain flexibility and generality!

---

## 10. Architecture Details

### Complete Transformer Encoder

$$Z = \text{LayerNorm}(X + \text{MultiHeadAttention}(X, X, X))$$
$$\text{Output} = \text{LayerNorm}(Z + \text{FFN}(Z))$$

**Component Breakdown:**

```
class TransformerEncoder(nn.Module):
    def __init__(self, d_model=512, num_heads=8, d_ff=2048, dropout=0.1):
        super().__init__()

        # Multi-head attention
        self.self_attn = MultiHeadAttention(d_model, num_heads)

        # Feed-forward network
        self.ffn = nn.Sequential(
            nn.Linear(d_model, d_ff),
```

```

        nn.ReLU(),
        nn.Dropout(dropout),
        nn.Linear(d_ff, d_model)
    )

    # Layer normalization
    self.norm1 = nn.LayerNorm(d_model)
    self.norm2 = nn.LayerNorm(d_model)

    # Dropout
    self.dropout1 = nn.Dropout(dropout)
    self.dropout2 = nn.Dropout(dropout)

    def forward(self, x, mask=None):
        # Self-attention with residual
        attn_output = self.self_attn(x, x, x, mask)
        x = self.norm1(x + self.dropout1(attn_output))

        # Feed-forward with residual
        ffn_output = self.ffn(x)
        x = self.norm2(x + self.dropout2(ffn_output))

    return x

```

## Complete Transformer Decoder

```

class TransformerDecoder(nn.Module):
    def __init__(self, d_model=512, num_heads=8, d_ff=2048, dropout=0.1):
        super().__init__()

        # Masked self-attention (causal)
        self.self_attn = MultiHeadAttention(d_model, num_heads)

        # Cross-attention to encoder
        self.cross_attn = MultiHeadAttention(d_model, num_heads)

        # Feed-forward
        self.ffn = nn.Sequential(
            nn.Linear(d_model, d_ff),
            nn.ReLU(),
            nn.Dropout(dropout),
            nn.Linear(d_ff, d_model)
        )

        # Layer norms
        self.norm1 = nn.LayerNorm(d_model)
        self.norm2 = nn.LayerNorm(d_model)
        self.norm3 = nn.LayerNorm(d_model)

        # Dropout
        self.dropout1 = nn.Dropout(dropout)
        self.dropout2 = nn.Dropout(dropout)
        self.dropout3 = nn.Dropout(dropout)

    def forward(self, x, encoder_output,
               src_mask=None, tgt_mask=None):
        # Masked self-attention
        attn1 = self.self_attn(x, x, x, tgt_mask)
        x = self.norm1(x + self.dropout1(attn1))

        # Cross-attention to encoder
        attn2 = self.cross_attn(x, encoder_output,
                               encoder_output, src_mask)
        x = self.norm2(x + self.dropout2(attn2))

        # Feed-forward
        ffn_out = self.ffn(x)
        x = self.norm3(x + self.dropout3(ffn_out))

```

```
    return x
```

---

## Training Details: Teacher Forcing

### During Training:

```
# Use ground truth as input (parallel training)
# Input: [SOS, "Le", "chat", "s'est"]
# Target: ["Le", "chat", "s'est", "assis", "EOS"]

for batch in dataloader:
    encoder_out = encoder(source)

    # Decoder sees ground truth (shifted right)
    decoder_out = decoder(target[:-1], encoder_out)

    # Predict next token
    loss = criterion(decoder_out, target[1:])

    loss.backward()
    optimizer.step()
```

### During Inference:

```
# Autoregressive generation (sequential)
generated = [SOS]

for _ in range(max_length):
    decoder_out = decoder(generated, encoder_out)
    next_token = argmax(decoder_out[-1])

    if next_token == EOS:
        break

    generated.append(next_token)
```

---

## Practical Hyperparameters

### Original Transformer ("Attention is All You Need"):

```
config = {
    'd_model': 512,           # Model dimension
    'num_layers': 6,          # Encoder and decoder layers
    'num_heads': 8,           # Attention heads
    'd_ff': 2048,            # FFN inner dimension
    'dropout': 0.1,           # Dropout rate
    'max_seq_len': 512,       # Maximum sequence length
    'vocab_size': 37000,      # Vocabulary size
}
```

### Transformer-Big (Better performance):

```
config_big = {
    'd_model': 1024,
    'num_layers': 6,
    'num_heads': 16,
    'd_ff': 4096,
    'dropout': 0.3,
}
```

### Modern Large Models (GPT-3, etc.):

```

config_large = {
    'd_model': 12288,           # 12K dimensions!
    'num_layers': 96,            # 96 layers
    'num_heads': 96,
    'd_ff': 49152,
    'params': '175B',          # 175 billion parameters
}

```

## Optimization: Learning Rate Schedule

**Warmup + Decay:**

```

def get_lr(step, d_model=512, warmup_steps=4000):
    arg1 = step ** (-0.5)
    arg2 = step * (warmup_steps ** (-1.5))
    return (d_model ** (-0.5)) * min(arg1, arg2)

```

$$\text{lr} = d_{model}^{-\frac{1}{2}} \cdot \min(\text{step}^{-\frac{1}{2}}, \text{step} \cdot (\text{warmup-steps}^{-\frac{3}{2}}))$$

$$\text{lr} = \frac{1}{\sqrt{d_{model}}} \cdot \min \left( \frac{1}{\sqrt{\text{step}}}, \text{step} \cdot \left( \frac{1}{\sqrt{\text{warmup-steps}^3}} \right) \right)$$

**Why Warmup?**

- Transformers sensitive to initialization
- Large gradients early in training
- Warmup stabilizes training

## 11. Computational Complexity Analysis

### Attention Complexity

**Standard Self-Attention:**

$$\text{Complexity} = O(n^2 \cdot d)$$

where:

- $n$ : sequence length
- $d$ : model dimension

**Breakdown:**

1.  $QK^T$ :  $O(n^2 \cdot d)$  — matrix multiplication
2. Softmax:  $O(n^2)$  — row-wise operation
3. Multiply by  $V$ :  $O(n^2 \cdot d)$

**Total:**  $O(n^2 \cdot d)$  per layer

**For full Transformer:**

- Encoder:  $L \cdot O(n^2 \cdot d)$  where  $L$  = number of layers
- Decoder: Similar but with additional cross-attention

### Memory Complexity

**Storing Attention Matrices:**

$$\text{Memory} = O(n^2 \cdot h + n \cdot d)$$

where  $h$  is number of heads

#### Components:

- Attention weights:  $h \times n \times n$  (can be large!)
- Activations:  $n \times d$  per layer
- Gradients: Same as activations (during training)

#### Example:

```
# Sequence length n = 1024, d_model = 512, heads = 8
attention_matrix = 8 * 1024 * 1024 * 4 bytes # ~33 MB per layer
# For 12 layers: ~400 MB just for attention!
```

---

## Comparison with Other Architectures

Architecture	Complexity per Layer	Sequential Ops	Max Path Length
RNN	$O(n \cdot d^2)$	$O(n)$	$O(n)$
CNN	$O(k \cdot n \cdot d^2)$	$O(1)$	$O(\log_k(n))$
Attention	$O(n^2 \cdot d)$	$O(1)$	$O(1)$

#### Trade-offs:

- RNN: Sequential bottleneck, but linear memory
- CNN: Limited receptive field, but efficient
- Transformer:  $O(n^2)$  complexity, but parallel + global

---

## Efficient Attention Variants

Problem:  $O(n^2)$  doesn't scale to very long sequences

#### Solutions:

##### 1. Sparse Attention (Routing patterns):

```
# Only attend to local neighbors + few global tokens
attention_pattern = local_window + strided + global_tokens
# Reduces to O(n * k) where k << n
```

##### 2. Linear Attention (Kernelized):

```
# Approximate attention with kernel functions
# O(n) complexity!
attention = φ(Q) @ (φ(K)^T @ V)
# instead of (Q @ K^T) @ V
```

##### 3. Flash Attention (I/O efficient):

```
# Fused CUDA kernels
# Same complexity but much faster in practice
# Reduces memory from O(n^2) to O(n)
```

---

## 12. Why Transformers Dominate

### Advantages Over Previous Architectures

Aspect	RNN/LSTM	CNN	Transformer
Parallelization	✗ Sequential	✓ Parallel	✓ Parallel
Long-range deps	✗ Gradient vanishing	✗ Limited receptive field	✓ Direct connections
Inductive bias	Strong (sequential)	Strong (locality)	Weak (flexible)
Data efficiency	✓ Good	✓ Good	✗ Needs large data
Interpretability	✗ Hidden state	✗ Feature maps	✓ Attention weights
Scalability	✗ Limited	✗ Saturates	✓ Continues improving

## The Scaling Hypothesis

**Key Observation:** Transformers continue to improve with:

- More parameters
- More data
- More compute

$$\text{Performance} = \text{Power Law}(Scale)$$

$$\log(\text{Loss}) \approx -\alpha * \log(N) + C$$

where N = parameters, data, or compute

**This doesn't hold for CNNs/RNNs!** They saturate.

## Empirical Wins

**NLP:**

- Machine Translation: BLEU score improvements
- Language Modeling: Perplexity reductions
- Few-shot Learning: Emergent capabilities

**Vision:**

- ImageNet: State-of-the-art accuracy
- Object Detection: Superior to CNN-based methods
- Video Understanding: Temporal modeling

**Multi-modal:**

- CLIP: Zero-shot classification
- DALL-E: Text-to-image generation
- Flamingo: Few-shot learning across modalities

## The Lottery: What Makes Transformers Special?

**Hypotheses:**

1. **Inductive Bias Trade-off:**
  - Weak priors allow learning from data
  - Doesn't impose wrong assumptions
2. **Expressiveness:**
  - Universal approximation with fewer constraints
  - Can represent more complex functions
3. **Optimisation Landscape:**
  - Residual connections create smooth paths

- Skip connections prevent gradient issues

#### 4. Attention as Routing:

- Soft, learnable connectivity
- Adaptive computation based on content

#### 5. Parallel Training:

- Efficient use of modern hardware
- Scales better with resources

## 13. Limitations and Challenges

### Computational Cost

```
# Example: GPT-3 training
parameters = 175e9
tokens = 300e9
compute = 3.14e23 FLOPs # ~$4.6M in cloud costs!
training_time = ~1 month on thousands of GPUs
```

#### Inference Cost:

```
# Single forward pass for GPT-3
sequence_length = 2048
floating_point_ops = 2 * 175e9 * 2048 # ~700 GFLOPs
# Hundreds of ms latency
```

## Quadratic Complexity in Sequence Length

#### Problem:

```
# Memory and compute grow as O(n2)
seq_len = [128, 256, 512, 1024, 2048, 4096]
memory = [x**2 for x in seq_len]
# [16K, 65K, 262K, 1M, 4M, 16M] elements

# Can't process very long sequences!
# Books, long documents, high-res images, videos
```

#### Partial Solutions:

- Sparse attention patterns
- Linear attention approximations
- Hierarchical processing
- Segmentation strategies

## Data Hunger

**Observation:** Transformers need massive amounts of data

```
# Typical requirements
small_model = {
    'params': '110M',
    'data': '10B tokens', # ~10GB text
    'compute': '1e20 FLOPs'
}

large_model = {
    'params': '175B',
    'data': '300B tokens', # ~300GB text
```

```
        'compute': '3e23 FLOPs'  
    }
```

## Why?

- Weak inductive biases
- Must learn structure from data
- Overparameterization requires regularization through data

---

## Interpretability Challenges

### Attention ≠ Explanation:

```
# Common misconception:  
# "High attention weight means the model uses this information"  
  
# Reality:  
# - Attention is just one component  
# - FFN can override attention  
# - Multiple heads complicate interpretation  
# - Attention can be uniform but still useful
```

### Open Questions:

- What do different heads learn?
- Why do some heads become "no-op"?
- How does information flow through layers?
- What concepts do neurons represent?
- how is knowledge actually stored?

---

## 14. Modern Variants and Extensions

### Encoder-Only Models

#### BERT (Bidirectional Encoder Representations from Transformers):

```
# Masked language modeling  
input: ["The", "[MASK]", "sat", "on", "the", "mat"]  
output: ["The", "cat", "sat", "on", "the", "mat"]  
  
# Used for:  
# - Classification  
# - Named Entity Recognition  
# - Question Answering
```

#### Architecture:

- Only encoder stack
- Bidirectional attention (see full context)
- Pre-trained on massive corpora

---

### Decoder-Only Models

#### GPT (Generative Pre-trained Transformer):

```
# Autoregressive language modeling  
input: ["The", "cat", "sat"]  
output: "on" # Predict next token  
  
# Used for:
```

```
# - Text generation  
# - Few-shot learning  
# - In-context learning
```

#### Architecture:

- Only decoder stack (with causal masking)
- Unidirectional attention
- Scales to massive sizes (GPT-3: 175B params)

**Key Insight:** Decoder-only models can do everything!

- Generation (natural)
- Classification (with prompting)
- Translation (with examples)

---

## Efficient Transformers

**Reformer** (Kitaev et al., 2020):

- Locality-sensitive hashing for attention
- Reversible layers (saves memory)
- $O(n \log n)$  complexity

**Linformer** (Wang et al., 2020):

- Low-rank approximation of attention
- $O(n)$  complexity
- Small accuracy drop

**Performer** (Choromanski et al., 2021):

- Kernelized attention (FAVOR+)
- $O(n)$  complexity
- Maintains performance

---

## 15. Summary & Key Takeaways

### Core Concepts Recap

#### 1. Attention Mechanism:

$$\text{Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

- Soft, differentiable routing of information
- Content-based, not position-based

#### 2. Transformer Architecture:

- Stacked encoder-decoder with attention
- Multi-head attention for diverse patterns
- Position-wise feed-forward networks
- Residual connections + Layer normalization

#### 3. Why It Works:

- $O(1)$  path between any two positions
- Parallel processing (efficient training)
- Flexible, learned inductive biases
- Scales with compute and data

## When to Use Transformers

### Use Transformers When:

- Large dataset available
- Long-range dependencies important
- Parallel training resources available
- Flexibility more important than efficiency
- State-of-the-art performance needed

### Consider Alternatives When:

- Small dataset (<10K examples)
- Real-time inference critical
- Memory/compute constrained
- Strong domain priors available
- Interpretability is paramount

---

## The Road Ahead

### Current Research Directions:

#### 1. Efficiency:

- Linear attention mechanisms
- Sparse transformers
- Model compression

#### 2. Long Context:

- Extending to 100K+ tokens
- Hierarchical processing
- Memory-augmented transformers

#### 3. Multimodal:

- Unified architectures
- Cross-modal learning
- Few-shot transfer

#### 4. Theory:

- Understanding what is learned
- Optimization dynamics
- Generalization bounds

---

## 16. Practical Tips for Implementation

### Debugging Checklist

```
# Common issues and solutions

1. **Exploding/Vanishing Gradients**
   - Check: Gradient norms
   - Fix: Gradient clipping, learning rate warmup

2. **Attention Collapse**
   - Check: Attention weight entropy
   - Fix: Dropout, smaller initialization

3. **Training Instability**
   - Check: Loss spikes
   - Fix: Mixed precision, gradient accumulation

4. **Poor Performance**
   - Check: Positional encoding, masking
   - Fix: Verify masks, try learned positions
```

- 
5. **\*\*OOM (Out of Memory)\*\***  
– Check: Batch size, sequence length  
– Fix: Gradient checkpointing, smaller batches

---

## Hyperparameter Tuning Guide

**Start Here (Defaults):**

```
config = {  
    'learning_rate': 1e-4,  
    'warmup_steps': 4000,  
    'batch_size': 32,  
    'dropout': 0.1,  
    'num_layers': 6,  
    'num_heads': 8,  
    'd_model': 512,  
    'd_ff': 2048,  
}
```

**Tuning Priority:**

1. Learning rate (most impactful)
2. Warmup steps
3. Batch size
4. Dropout
5. Architecture (last resort)

---

## Pre-training vs Fine-tuning

**Pre-training** (if you have massive compute):

```
# Train from scratch on large corpus  
# Requires: 100GB+ data, weeks of training, $$$  
  
model = Transformer(...)  
train(model, large_corpus, epochs=100)
```

**Fine-tuning** (recommended for most):

```
# Start from pre-trained model  
# Requires: Small dataset, hours of training, $  
  
model = load_pretrained('bert-base')  
model.classifier = nn.Linear(768, num_classes)  
  
# Lower learning rate for fine-tuning  
optimizer = AdamW(model.parameters(), lr=2e-5)  
train(model, task_data, epochs=3)
```

---

## 17. Connections to Future Topics

### What's Next in This Course

**Week 6: Sequence Models (RNNs, LSTMs, S4)**

- How Transformers evolved from RNNs
- Trade-offs between architectures
- When to use what

## Week 7: Self-Supervised Learning

- BERT and GPT pre-training objectives
- Contrastive learning with Transformers
- Masked modeling strategies

## Week 10: Scaling Laws

- Why Transformers scale so well
- Emergent abilities
- Data/compute trade-offs

## Week 12: Mechanistic Interpretability

- What Transformers learn
  - Attention patterns
  - Circuit discovery
- 

# 18. References & Further Reading

## Essential Papers

1. [Attention Is All You Need](#) (Vaswani et al., 2017)
    - The original Transformer paper
    - Must-read foundation
  2. [BERT](#) (Devlin et al., 2018)
    - Bidirectional pre-training
    - Encoder-only architecture
  3. [GPT-2](#) and [GPT-3](#) (OpenAI)
    - Decoder-only at scale
    - Few-shot learning emergence
  4. [Vision Transformer \(ViT\)](#) (Dosovitskiy et al., 2020)
    - Covered in Lecture 04
    - Transformers for images
  5. [Formal Algorithms for Transformers](#) (Phuong & Hutter, 2022)
    - Comprehensive mathematical treatment
    - Excellent reference
- 

## Tutorials and Code

- [The Annotated Transformer](#)
    - Line-by-line implementation
    - Best starting point for coding
  - [Hugging Face Transformers](#)
    - Production-ready implementations
    - Pre-trained models
  - [PyTorch Transformer Tutorial](#)
    - Official tutorial
    - Clean implementation
  - [minGPT](#)
    - Minimal GPT implementation
    - By Andrej Karpathy
- 

## Advanced Topics

- **Efficient Transformers:** [Survey Paper](#)

- **Scaling Laws:** [Kaplan et al.](#)
  - **Interpretability:** [A Mathematical Framework for Transformer Circuits](#)
  - **Optimization:** [On Layer Normalization](#)
- 

## Appendix: Mathematical Derivations

### A1: Why Scale by $\sqrt{d_k}$ ?

**Problem:** Without scaling, dot products grow large in high dimensions

**Proof:** Assume  $q_i, k_j \sim \mathcal{N}(0, 1)$  independently.

$$q \cdot k = \sum_{i=1}^{d_k} q_i k_i$$

$$[q \cdot k] = \sum_{i=1}^{d_k} [q_i][k_i] = 0$$

$$\text{ar}(q \cdot k) = \sum_{i=1}^{d_k} \text{ar}(q_i k_i) = d_k$$

So  $q \cdot k \sim \mathcal{N}(0, d_k)$

**Issue:** Softmax saturates when inputs are large!

$$f(x|0) : \text{softmax}(x) = [0, \dots, 1, \dots, 0]$$

(near one-hot  $\rightarrow$  small gradients)

**Solution:** Scale by  $\sqrt{d_k}$ :

$$\frac{q \cdot k}{\sqrt{d_k}} \sim \mathcal{N}(0, 1)$$

Now variance is constant regardless of dimension!

---

### A2: Positional Encoding Properties

**Claim:** For offset  $k$ ,  $PE_{pos+k}$  is a linear function of  $PE_{pos}$

**Proof sketch:**

Using trigonometric identities:

$$\sin(\alpha + \beta) = \sin(\alpha)\cos(\beta) + \cos(\alpha)\sin(\beta)$$

$$\cos(\alpha + \beta) = \cos(\alpha)\cos(\beta) - \sin(\alpha)\sin(\beta)$$

We can write:

$$[PE_{pos+k,2i} \ PE_{pos+k,2i+1}] = [\cos(\beta) \ \sin(\beta) \ -\sin(\beta) \ \cos(\beta)] [PE_{pos,2i} \ PE_{pos,2i+1}]$$

where  $\beta = \frac{k}{10000^{2i/d}}$

**Implication:** Relative position  $k$  encoded by fixed transformation! Model can learn to attend by relative position.

[DERIVATION NEEDED: Full mathematical proof]

---

### A3: Multi-Head Attention as Ensemble

**Theorem:** Multi-head attention with  $h$  heads learns  $h$  different attention patterns.

**Intuition:** Each head  $i$  has parameters  $(W_i^Q, W_i^K, W_i^V)$  that are learned independently.

The attention score for head  $i$ :

$$A_i = \text{softmax} \left( \frac{(XW_i^Q)(XW_i^K)^T}{\sqrt{d_k}} \right) (XW_i^V)$$

Final output combines all heads:

$$\text{Output} = \text{Concat}(A_1, \dots, A_h)W^O$$

Each  $A_i$  can specialize to different patterns!

**Empirical observation:**

- Some heads track syntax
- Some heads track semantics
- Some heads track positional patterns
- Redundancy provides robustness

---

## [END OF SLIDES]

---

## Next Steps for Students

### Before Next Class:

1. **Read:** "Attention Is All You Need" paper
2. **Code:** Implement scaled dot-product attention
3. **Experiment:** Train small Transformer on toy task
4. **Think:** How would you apply Transformers to your research area?

### Lab Assignment:

Implement a mini-Transformer for character-level language modeling:

- Dataset: Shakespeare text
- Task: Predict next character
- Model: 4-layer decoder-only Transformer
- Deliverable: Working model + analysis

### Project Ideas:

1. Compare Transformer vs RNN on long sequences
  2. Visualize attention patterns in trained model
  3. Implement efficient attention variant
  4. Apply Transformer to new domain (graph, time series, etc.)
- 
-